

Online Appendix

Trouble Downstream: How Rivers Carry Rainfall Shocks to Conflict in India

Mahiro Kobayashi[†]

October 7, 2026

[Click here for the latest online appendix](#) · [Main paper](#)

Contents

A	Data Construction	3
A.1	Unit of Analysis	3
A.2	Defining Upstream Catchment	3
A.3	The Treatment Weights	8
A.4	Measuring the Upstream Shock	11
B	Inference and Identification	18
B.1	Expected Exposure	18
B.2	The Number of Shocks	18
B.3	Standard Errors under Spatial Dependence	20
B.4	Randomization Inference	22
C	Routed Water: CaMa-Flood and Its Validation	24
C.1	The CaMa-Flood Model	24
C.2	Validation	26
C.3	Inundated Extent and the Flood Catalogue	27
C.4	The Flood Threshold	28

^{*}Graduate School of Public Policy, University of Tokyo.

[†]Last modified on October 7, 2026.

D Mechanism	29
D.1 The Timing of the Response	29
D.2 Agriculture, Prices and Labour	32
D.3 Public Employment	48
E Additional Figures and Tables	49
E.1 Descriptive Figures	49
E.2 Main-Result Robustness	53
E.3 The Dam Arm	66
E.4 The Flood Arm	67

Appendix A. Data Construction

Appendix A documents how the data and the treatment are constructed: the unit of analysis, the upstream catchment of each district, the weights on its sub-basins, and the rainfall measure.

A.1. Unit of Analysis

All series are harmonized to 2001-census district boundaries. The census map has 594 units: 593 districts, and the part of the former state of Jammu and Kashmir administered by Pakistan and China (Gilgit-Baltistan, Azad Jammu and Kashmir and the Shaksgam tract), which the map draws but the census did not enumerate. That territory holds 4.0 percent of UCDP events, mostly India–Pakistan fighting along the Line of Control, and is left out of the panel. Table E.3 shows that the headline estimates do not depend on leaving it out: they keep their sign and significance whether its events are reassigned to the nearest Indian district or it enters the panel as a district of its own. The district is the finest unit at which UCDP events can be used without loss: of the 89.7 percent of events UCDP locates to a district or better, 34.8 percent carry exact coordinates, 33.1 percent are placed within 25 km of a known point, and 32.1 percent are known only to lie in a district and are coded at its centroid. A finer unit — a grid cell or a river sub-basin — would have to drop that third of the events or assign it arbitrarily, and would misassign more of the 25 km events: placing each at random within its stated radius, the true location falls outside the assigned unit for 33 percent of them under districts and 67 percent under 0.25-degree cells.

A.2. Defining Upstream Catchment

The catchments are built from HydroBASINS (Lehner and Grill, 2013), a global map that divides land into sub-basins and records, for each, the one sub-basin its water flows into next, so that water can be followed from any sub-basin down to the sea (Figure A.1). The map is nested in twelve levels: each level splits every sub-basin of the level above into smaller ones, until they reach roughly 100–250 km², below which the map stops splitting (Table A.1). The paper builds catchments from level 10, and uses levels 4 and 5, whole river systems and their main branches, as clusters for inference (Section B).

Level 10 is fine enough to draw a catchment around each district, but no finer than rainfall is measured. At level 10 the median sub-basin covers 139 km², against 4,060 km² for the median district (about thirty sub-basins per district), and is 4.9 times a cell of the 0.05-degree CHIRPS rainfall grid; only 7.2 percent of sub-basins are smaller than one cell. A finer level would add almost nothing: level 11 has only 9 percent more sub-basins over India than level 10, and the same median size (136 against


Table A.1: The HydroBASINS hierarchy over India

Level	Number	Median km ²	Unit	Use in this paper
<i>HydroBASINS sub-basins</i>				
1	1	20,827,400	Continent (Asia)	
2	3	2,255,600	Continental regions	
3	9	748,800	Largest river basins	
4	34	75,500	Whole river systems	Inference clusters (Sec. B)
5	120	23,100	Main river branches	Inference clusters (Sec. B)
6	431	6,200		Rainfall coarsened (Fig. E.8)
7	1,486	1,600		Rainfall coarsened (Fig. E.8)
8	4,662	502		Rainfall coarsened (Fig. E.8)
9	12,393	204		Rainfall coarsened (Fig. E.8)
10	22,945	139		Treatment catchments
11	24,976	136		
12	25,029	136		
<i>For reference</i>				
	1	3,186,000	India	
	593	4,060	Districts (2001 census)	

Notes: HydroBASINS v1c (Lehner and Grill, 2013), Asia layer. Sub-basins are counted if they intersect India (2001 census districts); median areas are of those sub-basins, including any part outside India. Levels 1–3 were drawn by hand; from level 4 on, each sub-basin is divided into nine: its four largest tributaries and the five stretches of main river between them. A sub-basin is split only where two rivers each draining at least 100 km² meet, or where it exceeds 250 km², which is why the count stops growing after level 10. Reference rows: the area of India and of the median district on the 2001 census map, without the unenumerated Pakistan- and China-administered unit, on an equal-area projection; for India the area column gives its total area.

139 km²), since level-10 sub-basins are already close to the size at which HydroBASINS stops splitting.

(a) Level-10 sub-basins



(b) The Kosi–Ganga confluence, district boundaries in black




Figure A.1: The HydroBASINS level-10 delineation

Notes: Panel (a): level-10 sub-basin boundaries (blue) within India’s bounding box, national outline in black; boundaries are simplified for rendering at this scale. Panel (b): the Kosi–Ganga confluence at full resolution, with 2001-census district boundaries in black. Sub-basins are the units on which upstream rainfall is accumulated before it is assigned to districts through the drainage network.

A district’s upstream catchment is every level-10 sub-basin whose water reaches it, found by following the chain of downstream sub-basins until it enters the district. HydroBASINS records a single downstream sub-basin for each, so it cannot represent a river that divides.¹ The construction therefore adds the points where rivers split, taken from a hydrological routing model (Section A.3), and follows the water down every branch, each carrying its share of the flow.

Two kinds of sub-basin are then removed, so that the catchment holds only rain that reaches the district and that neither the district nor its neighbours received. The first drain to an inland lake rather than the sea, such as Sambhar Lake in Rajasthan or the closed lakes of southern Tibet: HydroBASINS links them to a nearby river only for bookkeeping, and their water reaches no district outside their own closed basin. The second are caught by the exclusion rule of Section III, which removes every sub-basin that overlaps the district or a bordering district. Across an international border the neighbours are the adjacent administrative units of Bangladesh, Nepal, Pakistan, China, Myanmar and Bhutan, at the level closest in size to an Indian district.

The result does not depend on how finely rainfall is measured. Figure E.8 keeps

¹HydroBASINS is built from a flow-direction grid in which every cell drains to exactly one neighbour. At a split, all of the river’s water therefore continues down one branch, whichever the elevation model favours, and the other branch appears as a new river that starts at or near the split and collects only the land draining directly into it. Below Farakka, for instance, the whole Ganga continues into Bangladesh as the Padma, while the Bhagirathi–Hooghly is recorded as a separate river whose catchment at Berhampore is a few thousand square kilometres, against nearly a million for the Ganga; in the Godavari delta the river continues down the Gautami branch, and the Vasishta starts afresh.

the catchment, exclusion and weights of Table 2 and coarsens only the rainfall: each level-10 sub-basin is given the average rainfall of the catchment’s sub-basins that lie in the same sub-basin at level 9, 8, 7 or 6 (Table A.1 gives their sizes). Even at level 6, where a sub-basin is larger than the median district, no coefficient moves by more than 0.16; columns 2 to 5 stay significant at the one-percent level throughout, and column 1, without the spatial controls, at the ten-percent level; column 4 is 1.344 against 1.460.

The exclusion has to sit between two failure modes, which a distance buffer makes visible (Section IV). Figure A.2 replaces the neighbour rule with a buffer, keeping only sub-basins more than a given distance from the district, and re-estimates the elasticity at every radius from 0 to 100 kilometres. With a narrow buffer the treatment is partly local weather: its correlation with the district’s own monsoon rainfall after fixed effects rises from 0.58 at 100 km to 0.91 with no buffer (panel e), and the fixed-sample elasticity from 0.57 to 1.61 (panel c). At both ends, the true district–catchment map no longer beats within-state permutations of it (panel f; Section B.4): the true estimate is 0.98 times the permutations’ 95th percentile at 100 km and 0.92 times with no buffer.

Three features of the sweep limit what it can show. First, the permutation test loses power as the radius widens. Districts in the same state increasingly share their remaining distant sub-basins — the weight two such districts have in common rises from 9 percent with no buffer to 19 percent at 100 km — so a permuted map comes closer to the true one, and the permutations’ mean rises from -0.08 to 0.26. Second, the eleven radii are not separate tests: each catchment is the one at the previous radius minus its nearest sub-basins, so neighbouring estimates share most of their data. Third, the districts that keep a catchment at 100 km are not a random subset: their catchments are larger (72 against 3.1 thousand km²), their sub-basins farther away (158 against 12 km) and their rivers flatter (0.91 against 0.56 flat share) — large floodplain districts whose water comes from far away, and not disproportionately in the Northeast.




Figure A.2: The lagged upstream elasticity across a continuous exclusion radius

Notes: Each radius keeps only upstream basins farther than that distance from the district. Panels (a)–(d) re-estimate columns 1 (Baseline) and 4 (Spatial controls) of Table 2, with Conley (150 km) 95 percent intervals; distance is measured to the basin polygon (a, c) or its representative point (b, d). In (a)–(b) the district set shrinks as the radius grows (districts retained along the bottom); (c)–(d) hold it at the districts that keep a catchment at 100 km. Panels (e)–(f) use the Baseline specification (Table 2, column 1), polygon distance and the varying district set, the blue series of (a): (e) is the correlation, after district and year effects, between the lagged treatment and the district’s own lagged monsoon rainfall; (f) compares the true estimate with 5000 permutations of the district–catchment map within states (filled points above the 95th percentile; exact two-sided p -values along the bottom). The paper’s own exclusion is not a distance rule, so no cell reproduces Table 2.

A.3. The Treatment Weights

The treatment weights each upstream sub-basin by how far its water travels along the river to reach the district: $\omega_{ij} = \sum_p s_p / (1 + r_p)$, summed over the paths p from sub-basin j to district i , where r_p is the path's length and s_p the share of j 's flow that follows it (Section III). Figure A.3 draws each piece of this notation. On a single channel there is one path with $s_p = 1$.

The length r_p runs along the river from the outlet of sub-basin j to where the water enters the district, the first sub-basin that overlaps it by more than 1 km². On a single channel it is the difference between the two sub-basins' distances to the river mouth, measured along the main channel and recorded in HydroBASINS. A path through a river split is measured piece by piece and summed: down the parent river to the split, along the branch channel at its length in CaMa-Flood, and down the branch, which HydroBASINS records as a river of its own (Section A.2), to the district.

The shares s_p come from the CaMa-Flood routing model, which gives, at each of the 85 points where a river divides, the share of the river's June–November flow that leaves through each branch, averaged over the estimation sample, 1989–2025. A path's share is the product of the shares of the branches it takes,

$$s_p = \prod_{k \in p} \phi_k,$$

where ϕ_k is the share of the flow arriving at a split that leaves through branch k . In Figure A.3 these are the labels beside the branches: the river leaving j first splits three ways, into a , b and $1 - a - b$, and the branch carrying b splits again, into c and $1 - c$. Path p_1 takes one branch, so $s_{p_1} = a$; path p_2 takes two, so $s_{p_2} = b \times c$, less than either of its branch shares. A path through no split has $s_p = 1$. A sub-basin that reaches the district by several paths adds them up, here $\omega_{ij} = a / (1 + r_{p_1}) + bc / (1 + r_{p_2})$, as in the figure's key. The branches marked $1 - a - b$ and $1 - c$ never reach i , so the path shares sum to less than one: $f_{ij} = a + bc < 1$. CaMa-Flood has no barrage or feeder canal at Farakka, whose design capacity is 4.3 percent of the modelled flow; setting the Farakka share to that capacity is one of the checks below.




Figure A.3: The notation of the treatment weights

Notes: Schematic, not to scale. Below the outlet of basin j (open circle) the river splits into three branches, with flow shares a , b and $1 - a - b$; the branch carrying b splits again, into c and $1 - c$. Two paths enter district i (filled circles; strictly, the first sub-basin overlapping i by more than 1 km^2), so $P(i, j) = \{p_1, p_2\}$. A path's share is the product of the branch shares it takes: $s_{p_1} = a$, $s_{p_2} = bc$. Its length r_p runs from the outlet to the entry point. The share of j 's water that reaches i is $f_{ij} = a + bc < 1$. The dashed straight line is the great-circle distance g_{ij} from i 's population-weighted centre (star) to j 's centroid, used by the alternative weights.

However large the catchment, much of the weight comes from within a few hundred kilometres of river. Figure A.4 shows the districts at the 90th percentile, the median and the 10th percentile of catchment size in the estimation sample: Dhubri (3,299 sub-basins; the district of Figure 2), Kokrajhar (217) and Malappuram (7). Half of their weight lies within 355, 172 and 154 km of the district along the river, although Dhubri's catchment reaches across Bhutan into Tibet.




Figure A.4: Upstream weights for three districts

Notes: Districts at the 90th percentile, median and 10th percentile of upstream sub-basins in the estimation sample. Maps shade each basin by Nw_{ij} , its weight relative to an equal share ($1 = \text{equal}$); grey basins are excluded (own, neighbouring or adjacent foreign units); the dashed outline holds the nearest basins carrying half the weight. White lines are rivers flowing into the district (HydroRIVERS, Strahler order ≥ 3), blue lines their course downstream. Panel (g): upstream basins per district.

The estimate does not depend on how the weights are built. Figure E.6 re-estimates columns 1 and 4 of Table 2 with the weights built in each of the six ways below, one

per row of the figure. In every case $w_{ij} = \omega_{ij} / \sum_{j' \in U(i)} \omega_{ij'}$; $f_{ij} = \sum_p s_p$ is the share of j 's water that reaches i , g_{ij} is the great-circle distance in kilometres from the district's population-weighted centre to the sub-basin's centroid, A_j is the sub-basin's area, and c_j is its marginal runoff per unit of monsoon rain, net of snowmelt.

1. *River distance (the paper's treatment)*: $\omega_{ij} = \sum_p s_p / (1 + r_p)$.
2. *Great-circle distance*: $\omega_{ij} = f_{ij} / (1 + g_{ij})$.
3. *Great-circle distance $\times$ basin area*: $\omega_{ij} = f_{ij} A_j / (1 + g_{ij})$.
4. *River distance $\times$ basin area*: $\omega_{ij} = \sum_p s_p A_j / (1 + r_p)$.
5. *Runoff volume: area $\times$ runoff coefficient*: $\omega_{ij} = f_{ij} A_j c_j$, with c_j measured over June–November.
6. *Runoff volume, decaying with river distance*: $\omega_{ij} = \sum_p s_p A_j c_j / (1 + r_p)$.

Across the five alternatives, column 1 stays between 0.83 and 0.95 and column 4 between 1.41 and 1.54.

Nor does it depend on the details of the construction. Three checks move neither column by more than 0.01: the population centre measured in 1988 (HYDE) rather than 2000, the runoff coefficient measured over June–September rather than June–November, and the Farakka share set to the feeder canal's design capacity. Figure E.7 varies how fast weight falls with distance, from equal weights to steep decay; column 4 stays between 1.31 and 1.50 throughout.

A.4. Measuring the Upstream Shock

The sign of the result does not depend on the rainfall product; its size does. The treatment is built from CHIRPS v3 (Funk et al., 2026), and Table A.2 sets it beside five alternatives built from gauges alone, a weather model, or an interpolated climatology. Each rebuilds the treatment from the same sub-basins, weights and monsoon months, so only the upstream rainfall changes. Figure A.5 re-estimates Table 2, column 4, under each product. The elasticity is positive under all six, on identical rows (panel (a)) and on each product's full coverage (panel (b)). On the common sample it is 1.50 on CHIRPS v3 and 1.02 on CHIRPS v2, against 0.26, 0.33, 0.24 and 0.26 on GPCC, CRU TS, ERA5-Land and TerraClimate.

Table A.2: The rainfall products

Product	Built from	Resolution	Record	Source
CHIRPS v3 (paper)	Satellite and gauges	0.05°	1981–present	Funk et al. (2026)
CHIRPS v2	Satellite and gauges	0.05°	1981–present	Funk et al. (2015)
GPCC v2022	Gauges only	0.25°	1891–2020	Schneider et al. (2014)
CRU TS 4.08	Gauges only	0.5°	1901–2023	Harris et al. (2020)
ERA5-Land	Reanalysis	0.1°	1950–present	Muñoz-Sabater et al. (2021)
TerraClimate	Interpolated climatology	1/24°	1958–present	Abatzoglou et al. (2018)

Notes: Record is the product’s own coverage. The estimates use each product within 1990–2025 as far as the files used here reach (Figure A.5, panel (b)).

The sizes differ because the products disagree about the year-to-year variation the estimate uses. Panel (c) of Figure A.5 regresses each product’s upstream rainfall on CHIRPS v3’s within district: the products move 0.51 to 0.77 log points per log point of CHIRPS v3, and CHIRPS v3 accounts for only 0.35 to 0.43 of their within-district variation.




Figure A.5: The lagged upstream elasticity across precipitation products

Notes: Each point re-estimates Table 2, column 4 — Poisson, district and year fixed effects, upstream temperature, local weather and the spatial block — replacing only the upstream rainfall with the product shown; local rainfall and the spatial block remain CHIRPS v3. Panel (a) uses the 6,324 district-years (204 districts, 1990–2020) that every product covers; panel (b) uses every district-year each product covers within 1990–2025 (years at right), so its CHIRPS v3 point is Table 2, column 4. In panels (a) and (b), bars are 95 percent intervals on Conley spatial-HAC standard errors (150 km) and the dashed line marks the CHIRPS v3 estimate. Panel (c) regresses each product's lagged log upstream rainfall on CHIRPS v3's with district and year fixed effects, on the common sample; bars are 95 percent intervals clustered by district, the grey dashed line marks a one-for-one relationship, and within R^2 is the share of the product's within-district variation that CHIRPS v3 accounts for.

CHIRPS v3 is the product best suited to this design, for three reasons. It combines satellite observation, which reaches the foreign headwaters, with gauges where they exist, whereas the gauge-only products interpolate across sparse Himalayan networks and India's gridded product (IMD) stops at the border, beyond which 24 percent of upstream sub-basins lie. Its 0.05-degree grid matches catchments built from sub-basins of about 139 km². And it covers the paper's whole window.

The magnitude should be read as a CHIRPS v3 magnitude, since version 3 differs substantially from version 2. Version 3 blends gauge data from nearly four times as many sources, corrects for gauge undercatch, uses a new 1991–2020 climatology,

and replaces the satellite algorithm, which in version 2 understated the year-to-year variance of rainfall and the size of heavy storms (Funk et al., 2026). The two versions correlate only 0.59 after district and year effects, and version 2 gives the smaller coefficient ($p = 0.010$). Either version 3 resolves rainfall the other products miss, or the difference is noise that happens to correlate with conflict; this paper cannot separate the two.

Over these catchments CHIRPS rests on few gauges, and most of them are foreign. The product reports the gauges it blended in each pixel and month (Figures A.6 and A.7). Few are Indian, since the IMD network does not flow into the product: 96 percent of Indian upstream sub-basins are effectively ungauged, against 83 percent of foreign ones, and the median estimation district draws 96 percent of its upstream weight from pixels with no station. The product blends about 82 gauges a monsoon month across India, and in the foreign catchments the count falls from 129 a month in the 1990s to 23 since 2011. Over these catchments CHIRPS is therefore close to a satellite product anchored by Nepal’s gauges.

Where gauge-based grids exist, CHIRPS agrees with them within districts, and least along the Himalayan arc (Figure A.8). District totals correlate 0.66 with IMD within district (median gap 14 percent), and sub-basin totals correlate 0.70 with GPCC in Indian and 0.61 in foreign sub-basins (gaps 11 and 22 percent).




Figure A.6: The gauges behind the treatment

Notes: Panel A maps CHIRPS v3’s own station-density diagnostic — the number of gauges the product blended, per 0.25-degree pixel — averaged over monsoon months (June–September) 1990–2025, masked to the upstream basins of the excl.-neighbor crosswalk. Red marks pixels averaging fewer than one blended station per twenty pixel-months. Panel B aggregates the same density to districts with the river-distance weights that build the upstream treatment, so it reads as the gauge support of each district’s exposure; quantile bins, darker better gauged. Grey districts have no upstream catchment in the construction.




Figure A.7: The number of gauges behind CHIRPS v3, by year and month

Notes: Each cell is the number of rain gauges CHIRPS v3 blended in that month, from the product's station-density diagnostic (per 0.25-degree pixel, summed over the area; a pixel on the boundary counts by the share of it inside). Panel (a) covers India; panel (b) the upstream catchments of the paper's districts that lie outside India, in Nepal, China, Bhutan and Bangladesh. Both panels share one colour scale. The box marks the monsoon months (June–September) of the paper's window, 1990–2025.

A. Districts: CHIRPS v3 vs IMD (India only)

B. Upstream basins: CHIRPS v3 vs GPCC




Figure A.8: CHIRPS v3 against the gauge-based grids

Notes: Panel A: the district's mean June–September CHIRPS v3 total (the panel's own local-rainfall measure) against the same total from the IMD 0.25-degree gauge-based grid, 1990–2025, as a percentage of the IMD total; blue means CHIRPS reads wetter. IMD covers India only, so the comparison cannot reach the foreign upstream basins. Panel B: the same comparison at basin level against GPCC v2022, the gauge-based product with global coverage, over the upstream basins of the treatment, 1990–2020. Gaps beyond ± 30 percent are censored to the scale's ends.

Poor measurement does not drive the result. Table A.3 splits districts into terciles of the gauge support of their treatment. The least-gauged tercile carries 1.550, significant at one percent, against 1.499 in the best-gauged, significant at ten percent, and the difference is not detectable ($p = 0.13$). Classical measurement error would attenuate the estimate where gauges are scarce, so a large estimate in the least-gauged tercile argues against a measurement artifact.

Dropping sub-basins is a weaker test, and its answers are mixed. Indian sub-basins alone give 0.82 (standard error 0.43, 155 districts), no longer significant, while dropping sub-basins above 2,500 m gives 0.86 and above 1,500 m gives 1.45 (standard error 0.35), larger than the headline. The result does not live in the badly measured high catchment, but it is sensitive to which foreign sub-basins enter. Such restrictions also change the estimand as well as the measurement, since snowmelt-fed sub-basins deliver water on a different schedule, and dropping the least-gauged tercile keeps a positive but imprecise 1.119 (standard error 0.646): they answer a power question, which is why the tercile split carries the test.

Table A.3: The headline estimate by gauge support of the treatment

	Full sample	Drop least- gauged third	Least gauged (T1)	Middle (T2)	Best gauged (T3)
log Upstream Precip. ($t - 1$)	1.460*** (0.283)	1.119* (0.646)	1.550*** (0.538)	0.611 (0.717)	1.499* (0.820)
Observations	7,560	4,824	2,736	2,592	2,232
Number of Districts	210	134	76	72	62

Notes: Poisson fixed-effects on the specification of Table 2, column 4, 1990–2025. Districts are split into terciles of the gauge support of the treatment: CHIRPS v3’s own monthly station counts (Figure A.6), averaged over each district’s upstream basins with the river-distance weights that build the exposure. The formal test is the interaction, which lets the treatment slope differ by tercile while holding every other coefficient common: the least-gauged tercile differs from the best-gauged by 1.05 ($p = 0.13$) and the middle by 0.80 ($p = 0.29$) — no heterogeneity by gauge support is detectable. The tercile columns, in which every coefficient may differ, run on 72–82 districts each and are shown for transparency; their movement is what samples of that size produce. Conley spatial-HAC standard errors (150 km cutoff) in parentheses, which is the paper’s headline estimator and the basis of the significance stars; Online Appendix Figure B.3 re-prices the headline coefficient under alternative variance estimators. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.

Appendix B. Inference and Identification

This appendix checks identification first and inference second. Section B.1 tests whether exposure to the shocks, rather than the shocks themselves, drives the estimate. The remaining sections address inference. Districts along the same river share upstream sub-basins, so their treatments are correlated by construction: rain over Nepal alone enters the treatment of 57 districts and is at least half of it for 26 (Figure B.1). Standard errors that treat districts as independent then overstate precision, and the number of independent shocks behind the estimate is far smaller than the number of districts (Adão et al., 2019). I address this in three ways. Section B.2 counts the effective number of shocks and checks that no single district or state drives the estimate. Section B.3 lets the errors correlate across districts up to whole river systems, under spatial-HAC kernels and under clusters of districts that share upstream water. Section B.4 drops the asymptotic approximation and runs randomization tests that reassign the district–catchment map and the year labels.

B.1. Expected Exposure

Rainfall shocks can be as good as random while exposure to them is not. A district with a large, wet catchment receives more upstream rain in every year, and if such districts differ in conflict for other reasons, the treatment is confounded. Borusyak and Hull (2023) show that this confounder is expected exposure μ_i , the treatment a district would receive on average over draws of the shocks. I address it by recentering: the treatment is replaced by its deviation from μ_i , estimated as the district’s mean treatment over the sample years. When climatology is stable, μ_i is constant over time and district effects already absorb it, so recentering cannot move the estimate. It reproduces the baseline 1.460 (1.460), which checks the construction rather than the assumption. The real test lets climatology drift by estimating μ_i separately for each decade. Expected exposure then carries 3.0 percent of the treatment’s variance, and the estimate falls to 1.390 (standard error 0.393, $p = 0.000$), still significant at one percent (Table B.1, panel B).

B.2. The Number of Shocks

Shift–share estimates are consistent as the number of independent shocks grows, not the number of districts (Borusyak et al., 2022), so what matters is how many shocks the treatment effectively draws on. I measure this as the inverse Herfindahl index $1/\sum_b \bar{s}_b^2$, the effective sample size that Borusyak et al. (2022, p. 204) recommend reporting. Here $\bar{s}_b$ is basin b ’s weight in the treatment averaged over districts (mapped in Figure B.2), and the index equals the number of basins when all weigh equally. It is 2,087 across 18,109 level-10 sub-basins, 37.4 across 82 major sub-catchments and

7.20 across 21 macro basins. Sub-basins of one river system share the same monsoon, so their shocks are not independent. Pooling them, as in the clustered version of the condition (Borusyak et al., 2022, Assumption 5), makes the macro-basin count the relevant one. The standard errors of Section B.3 therefore allow correlation up to that level.

No single district or state drives the estimate. Re-estimated with one district removed, it stays between 1.37 and 1.55; with one state removed, between 1.32 and 1.57, against 1.46 in the full sample (Figure E.13). Table B.1, panel A, collects the shock counts.




Figure B.1: Where each region's rain enters the treatment

Notes: Each panel takes one source region (outlined) and shades every district i by the share of its upstream-rainfall weight that falls in that region, $\sum_{j \in R} w_{ij}$. Basins are assigned to a region by their representative point: an Indian state, or the country outside India. Grey districts have no upstream catchment under the exclusion rule; white districts draw nothing from the region. Districts inside a source region can be white or pale because the exclusion rule drops their own and their neighbours' basins.




Figure B.2: Each basin's share of total exposure

Notes: Each basin b is shaded by $s_b = \frac{1}{N} \sum_i w_{ib}$, its exposure share summed over the 425 districts with an upstream catchment, relative to an equal split over the 22,237 basins that enter any district's treatment ($\sum_b s_b = 1$; 1 is an equal share). Weights as in Figure A.4.

B.3. Standard Errors under Spatial Dependence

Districts that share upstream water receive correlated treatments, and nearby districts may share unobserved shocks as well. If their errors are correlated, standard errors clustered by district are too small. I re-estimate the standard error of the headline coefficient (Table 2, column 4, 1990–2025) under eight alternatives, each allowing a different form of dependence:

- *Clustering by state*: arbitrary correlation among districts of the same state, which share administration and policing.
- *Clustering by macro basin*, at PFAF levels 4 and 5: arbitrary correlation among districts of the same river system, however far apart.
- *Clustering by catchment group*: arbitrary correlation among the 85 groups of districts whose upstream-weight vectors have cosine similarity above 0.50 (Figure E.10).

Table B.1: Shift-share design diagnostics

<i>A. How many independent shocks?</i>		
	Units	Effective number
HydroBASINS level-10 sub-basins	18,109	2,087
PFAF level-5 sub-catchments	82	37.4
PFAF level-4 macro basins	21	7.20
largest single macro-basin share	—	0.32
<i>B. Recentering on expected exposure</i>		
	Coefficient	Std. error
Baseline	1.460***	(0.235)
Recentered, time-invariant μ_i	1.460***	(0.235)
Recentered, μ_i by decade	1.390***	(0.393)
μ_i share of exposure variance	3.0%	—

Notes: All panels are computed on the estimation sample of Table 2, column 4 — 7,560 district-years over 210 districts, 1990–2025. Panel A gives $1/\sum_b \bar{s}_b^2$ over average exposure shares: the number of *independent* shocks the design has at each level of the basin hierarchy, which is what shift-share asymptotics run in. The sub-basin count should not be believed — sub-basins inside one monsoon system are not independent draws — which is why inference is clustered at the macro basin. Panel B recenters the treatment on expected exposure μ_i . With μ_i time-invariant, district effects absorb it exactly and the estimate cannot move — reproducing the baseline is a check on the construction, not a result. Letting μ_i differ by decade is the demanding version. Standard errors clustered by district. ** 0.05, *** 0.01.

- *Conley spatial HAC*, at cutoffs from 25 to 500 kilometres: correlation that decays with geographic distance, regardless of state or basin boundaries.
- *Network HAC*, with cosine or Jaccard kernels: correlation that rises with the upstream basins two districts share, regardless of distance.
- *Exposure-robust* standard errors at PFAF level 5 (Adão et al., 2019): the basin shocks rather than the districts are the sampling unit, the inference counterpart of the shock count in Section B.2.

The estimate clears one percent under every one of them: the smallest t -statistic is 4.97, against a critical value of 2.576 (Figure B.3).




Figure B.3: The headline test statistic under alternative variance estimators

Notes: Every point is the same regression — the specification of Table 2, column 4, estimated on 1990–2025 — re-priced under a different variance estimator, so the figure varies the estimator and nothing else. The point estimate is 1.46 throughout and t is that estimate divided by the standard error in question. The connected curve gives Conley spatial-HAC statistics at distance cutoffs from 25 to 500 kilometres; the grid stops at 500 because beyond it the kernel encompasses most of the country and no longer describes local dependence. The curve is not monotone in the cutoff, and at some cutoffs the variance matrix must be repaired before use; its level is informative and its wiggles are not. The eight triangles at the left have no distance cutoff and are shown at their own labelled positions. Five cluster the residuals — on districts, states, PFAF level-4 and level-5 macro basins, and the catchment grouping built from the treatment itself in this section. The last three do not cluster at all: ‘HAC cos’ and ‘HAC jac’ are spatial-HAC estimators whose kernel is the cosine, respectively Jaccard, similarity of districts’ upstream-weight vectors, so two districts fed by the same basins are treated as dependent however far apart they lie; ‘Expos. P5’ is the exposure-robust estimator of Adão et al. (2019), which treats the PFAF level-5 basin shocks rather than the districts as the sampling unit. The exposure-robust estimator computed at the finest basin level is deliberately omitted: with more than a thousand basin shocks its asymptotics do not bite and it returns a standard error of 0.060, far tighter than anything else here, which would imply a precision the design does not have. Dashed lines mark the two-sided normal critical values at the 10, 5 and 1 percent levels.

B.4. Randomization Inference

Randomization inference asks whether the estimate could arise from the structure of the data alone. Each test re-estimates the same regression after reassigning one part of the data at random, holding the rest fixed, and compares the actual estimate with the distribution of the reassigned ones; the exact p -value is the share of draws at least as extreme. The paper reports four such tests, each with 5,000 draws.

- *The district–catchment map, within states* (Figure E.11; Figure E.12, panel (a)). *Randomized:* which catchment’s rainfall series each district receives, among the districts of its state. *Held fixed:* each catchment’s full time series, each district’s conflict and local weather, and the panel. *Tests:* whether a district’s own catchment predicts its conflict better than a plausible neighbour’s. No permutation reaches the actual estimate of 1.46 ($p = 0.0002$); the draws centre on -0.087, with a 95th percentile of 0.31.

- *The year labels* (Figure E.12, panels (b) and (d)). *Randomized*: the years of the weather block — upstream rainfall and temperature and the district’s own weather — with one permutation applied to every district. *Held fixed*: the map, the conflict series and the spatial terms built from it, and each year’s cross-section of shocks, so that districts sharing sub-basins keep their correlated shocks. *Tests*: whether conflict follows the timing of the shocks. The actual estimate lies outside the null on the coefficient ($p = 0.0040$), the studentized statistic ($p = 0.0002$, against an actual t of 6.22) and the score statistic implied by the Poisson likelihood ($p = 0.0080$).
- *The year labels, within decades* (Figure E.12, panel (c)). As the previous test, but each year moves only within its decade, which preserves trends and relaxes the exchangeability the free permutation assumes. *Tests*: the same question without relying on trends being exchangeable ($p = 0.0002$ studentized, $p = 0.0008$ score).
- *The district–catchment map at each exclusion radius* (Figure A.2, panel (f)). *Randomized*: as in the first test, at each radius from 0 to 100 kilometres, under the Baseline specification. *Held fixed*: as in the first test. *Tests*: at which buffer widths the true map still beats a plausible neighbour’s (Section A.2).

The two kinds of test answer different questions, and neither subsumes the other. The map tests shuffle space and hold time; the year-label tests shuffle time and hold space. Permuting the map within states makes it demanding: a counterfeit map still hands each district rainfall correlated with its own, so the test asks not whether upstream rainfall matters but whether *this* district’s catchment matters more than a plausible neighbour’s.

Appendix C. Routed Water: CaMa-Flood and Its Validation

C.1. The CaMa-Flood Model

Upstream rainfall is a proxy for the water that reaches the district. Its weights (Section A.3) are statistical choices, not hydrology, and the same monsoon can flood a district or not depending on how wet the soil already is, how much falls as snow, and how the river network carries and stores the water (Figure C.1). To measure what the river delivers, I route runoff — rain that the land surface has already passed through soil, snow and evaporation — through CaMa-Flood (Yamazaki et al., 2011, 2013), a global river routing model, at six-arc-minute resolution over South Asia for 1989–2025.




Figure C.1: From monsoon rain to flood-stage days

Notes: Top row: the routed measure. ERA5-Land turns weather into soil moisture, snowpack and evaporation (1) and the runoff that leaves the land each day (2); CaMa-Flood delivers that runoff to the river channels (3) and routes it downstream, storing it in channels and spilling it onto floodplains, to give daily discharge at the district’s river cell, from which flood-stage days are counted (4). Weather is the only source of year-to-year variation; both models are fixed physics on fixed geography, and no discharge observations are assimilated. Bottom row: the treatment of Section III, a weighted average of sub-basin rainfall that passes through neither the land surface nor the river.

CaMa-Flood moves water along a fixed river network. The land is divided into unit catchments, each draining into the next one downstream, and each holds a river channel with a floodplain beside it, described by the channel’s width, bank height and length and by a floodplain elevation profile taken from high-resolution topography. Every time step, runoff enters the channels, the water stored in each unit catchment is updated by mass balance, and outflow to the next catchment is computed from the slope of the water surface between the two. When storage exceeds the channel’s capacity, water spills onto the floodplain, which slows and spreads the flood wave instead of passing it downstream unchanged.

Formally, the storage V_c (m^3) of unit catchment c evolves by

$$V_c^{t+\Delta t} = V_c^t + \sum_{k \in u(c)} Q_k^t \Delta t - Q_c^t \Delta t + A_c R_c^t \Delta t,$$

where $u(c)$ are the unit catchments draining directly into c , Q is outflow ($\text{m}^3 \text{s}^{-1}$), A_c is the catchment's area and R_c its runoff (m s^{-1}) (Yamazaki et al., 2011, eq. 9). Storage up to the channel's capacity stays in the channel; beyond it, the floodplain profile splits storage between channel and floodplain at a common water level, which gives the water depth and the flooded area (Yamazaki et al., 2011, eqs. 1–2). Outflow follows the local inertial approximation to the shallow-water equations (Bates et al., 2010),

$$Q_c^{t+\Delta t} = \frac{Q_c^t + \Delta t g A S_c}{1 + \Delta t g n^2 |Q_c^t| / (H^{4/3} A)},$$

where S_c is the water-surface slope to the downstream catchment, A the flow cross-section, H the hydraulic radius, n Manning's roughness and g gravity (Yamazaki et al., 2013, eq. 6). Where a river divides, bifurcation channels link unit catchments that the river network does not connect, and water flows through them toward whichever side has the lower water surface (Yamazaki et al., 2014b); $u(c)$ then includes catchments linked to c through them. These are the river splits of Section A.3.

The configuration is the model's South Asia set-up. The river map is built from MERIT Hydro (Yamazaki et al., 2019); channel widths come from satellite observation (Yamazaki et al., 2014a), filled by an empirical relation with discharge where none exists, and bank heights from that relation. Manning's n is 0.03 in channels and 0.10 on floodplains, the water level at river mouths is fixed, and there are no reservoirs or levees except in the dam-operated variant of Section VII. Runoff is daily ERA5-Land (Muñoz-Sabater et al., 2021), distributed to unit catchments by area. The simulation runs in three chunks (1988–1999, 2000–2018 and 2018–2025), each started from empty storage with a one-year spin-up that is discarded; re-runs from stored states reproduce the baseline bit for bit.

The measure keeps only water that arrives from upstream. For district d , a masked run removes the runoff of every unit catchment within 50 kilometres of the district,

$$\tilde{q}_{kt}^{(d)} = A_k R_{kt} \mathbf{1}[k \notin M_{g(d)}],$$

where $M_{g(d)}$ is the union of the 50-kilometre masks of the districts that share d 's run: districts whose masks barely overlap are simulated together, in 34 runs, each district losing at most 5 percent of its upstream input to the others' masks. Simulated discharge at river cell c is then a deterministic, nonlinear function of the history of

masked runoff upstream of c and of Θ , the time-invariant river and floodplain geometry and parameters of the configuration above,

$$Q_{ct}^{(d)} = \mathcal{F}_c\left(\{\tilde{q}_{ks}^{(d)}\}_{k \in \mathcal{U}(c), s \leq t}; \Theta\right),$$

where $\mathcal{U}(c)$ is every catchment upstream of c , including those reaching it through bifurcation channels. With local runoff masked, upstream runoff is its only time-varying input. A district’s flood-stage days in year y count the monsoon days on which discharge at its river cell with the largest upstream area, c_d , exceeds that cell’s two-year return level,

$$\text{FloodDays}_{dy} = \sum_{t \in \text{June--September of } y} \mathbf{1}\left[\bar{Q}_{cdt}^{(d)} > Q_{cd}^{\text{RP2}}\right],$$

where $\bar{Q}$ is the centred three-day mean of daily outflow and Q_c^{RP2} is a Gumbel fit, by L-moments, to the 37 annual maxima of the unmasked baseline at cell c , 1989–2025. Upstream rainfall generates these days mechanically: a log-point of upstream monsoon rain adds 2.5 flood-stage days (standard error 0.26). Section VI reports the race between the rainfall total and these routed days (Table 4).

The mask’s footprint does not drive the measure. Rebuilding the ensemble with the mask drawn at the paper’s neighbour-exclusion footprint instead of 50 km gives instruments correlating 0.94 within district and a flood-stage result of 0.024 (50 km) against 0.021 (neighbour exclusion) on the common sample. The 50-kilometre construction is the better identified of the two: excluding districts whose mask removes under 1 percent of upstream input leaves its estimate at 0.020 while the administrative variant loses significance, and the one neighbour-exclusion run that failed its volume-conservation gate did so for exactly this reason and is held out.

C.2. Validation

Against an independent routing model.—Against VIC-RRM — an independent model with different physics and routing, calibrated on Bangladesh gauge observations — CaMa’s daily discharge agrees closely: $r = 0.96$ at the Ganges checkpoint, median $r = 0.90$ across 365 district gauge cells, and median $\kappa = 0.48$ on the margin the design actually uses, which days exceed a cell’s own return level. The models share ERA5-Land forcing, so this validates routing rather than meteorology. The resulting ordering is respected throughout: discharge is validated twice over, inundated extent is biased in level but valid in rank, and modelled flood *duration* is not used.

Validation against satellite flooding.—Over the 147 catalogued GFD events, the model’s inundation extent correlates with observed flooded fractions at 0.41 (Spear-

man, district level), reaching 0.51 with a probability of detection of 0.90 at unit-catchment resolution on major events. The model is naturalized — no reservoirs or levees — so infrastructure-failure floods like the 2008 Kosi embankment breach are invisible to it by construction: the routed instrument captures rain-driven floods only, precisely the exogeneity property the design wants.




Figure C.2: Routed and observed flooding over time

Notes: Yearly means across the districts carrying both sources (387 districts with a CaMa river cell, an upstream signal and no bifurcation flag). Panel (a) is in levels — the monsoon-maximum percentage of district area flooded — for GFD (satellite-observed, any origin), the CaMa baseline run (modelled, any origin), and the CaMa masked run (modelled, upstream origin only). Panel (b) is the paper’s treatment: days the three-day-mean upstream-origin discharge exceeds the two-year return level, averaged over the same districts. GFD district-years lost to cloud cover, and districts without CaMa floodplain cells, drop from their series’ means.

C.3. Inundated Extent and the Flood Catalogue

Extent and depth.—Inundated *area* and flood *depth* — the margins an optical flood map measures — carry no effect and absorb none of the flood-stage coefficient. This is not an artifact of measurement scale: rebuilt like-for-like, as days on which the district’s flooded fraction exceeded its *own* return level, the footprint remains a tight zero, alone and raced against discharge on identical observations. Depth so normalized does move conflict, but depth-days and discharge-days correlate at 0.94 within district — two readings of the same river — so it corroborates rather than competes. Nor is weighting the explanation: rebuilt weighting each pixel by population or kharif harvested area, the footprint still carries nothing.

The footprint measures are also validated on their own terms: on the Dartmouth catalogue’s impact fields across 144 usable India events, a log point of flooded area raises recorded deaths by 0.32 and displacement by 0.73. The measures that best predict who dies predict nothing about conflict, while routed discharge days continue to — humanitarian severity and conflict relevance are different quantities.

The footprint’s null is adaptation, not irrelevance: under abnormal inundation land moves *into* flood-tolerant rice and *out* of waterlogging-sensitive cash crops — the flood-tolerance margin of [Emerick et al. \(2016\)](#). Spread water is absorbed by

land use; the channel-force margin, dominated by the first days at flood stage, is the margin the portfolio cannot insure, and the conflict response concentrates in the paddy-floodplain systems whose agriculture engages river water, not where water-vulnerable crops dominate. One methodological note: the flood channel operates chiefly within the year, the agricultural-damage channel into the following one, so testing the flood channel against the wrong lag will fail.

Modeled flooding versus the flood catalogue.—The natural alternative to a routing model is the Global Flood Database (Tellman et al., 2021), and the reason to decline it is measurement (Patel, 2024). Every one of its 147 India events descends from the news-based Dartmouth catalogue, and the error is testable in-sample: small floods are missing non-randomly, detection depends on who lives downstream, and detection varies with cloud cover — which tracks the storm itself. The regression the catalogue invites shows the consequence: a catalogued-event indicator predicts *fewer* recorded conflict events, but the GFD’s own mapped extent carries nothing, a CaMa-built indicator matched to the catalogue’s base rate is significantly *positive*, and the negative appears identically in two independent news-derived conflict datasets — attention displacement within a shared news pipeline, not pacification. The design therefore uses continuous physical quantities throughout and retains the GFD in the validation role below.

C.4. The Flood Threshold

Does the flood threshold define the result?—The return level is what makes a day a flood day. Refitting the levels on the full 37-year routing record rather than the 19 annual maxima the published panel used raises the two-year level by a median 4.5% and cuts mean flood-stage days on the race sample from 2.22 to 1.87, while the two series still correlate at 0.973. Put side by side in the specification that races them, nothing moves: routed flood days read 0.0277 against 0.0276, upstream rainfall 0.781 against 0.788. The ordering Section VI draws on does not depend on where the threshold was placed.

Appendix D. Mechanism

This appendix collects the evidence on how the upstream shock reaches conflict: when the response appears (Appendix D.1), what happens to agriculture, prices and labour (Appendix D.2), and what public employment does (Appendix D.3).

D.1. The Timing of the Response

The main text dates the treatment to a year; this subsection asks when within the following two years the response appears, and whether it is specific to the monsoon or merely to rainfall.

Re-indexing.—A district-month is relabelled by monsoon year v (June to May) and position $k \in \{0, \dots, 11\}$, months since June — a complete relabelling of the calendar, with $y = v + \mathbb{1}\{k \geq 7\}$. Let C_{ivk} count events UCDP places in district i in that month, located to a district and dated to a week or better.

Specification.—Let $R_{i,v}^M$ be upstream monsoon rainfall — exactly the main text’s treatment — and $R_{i,v}^N$ the complementary January–May and October–December total, both one number per district-year. With T^M and T^N the matching upstream temperatures,

$$\begin{aligned} \log \mathbb{E}[C_{ivk}] = & \beta_k^M \log R_{i,v}^M + \beta_{k+12}^M \log R_{i,v-1}^M + \beta_k^N \log R_{i,v}^N + \beta_{k+12}^N \log R_{i,v-1}^N \\ & + \gamma_k^M T_{i,v}^M + \gamma_{k+12}^M T_{i,v-1}^M + \gamma_k^N T_{i,v}^N + \gamma_{k+12}^N T_{i,v-1}^N \\ & + \mathbf{x}_{ivk}' \boldsymbol{\theta} + (\mathbf{W} \mathbf{z}_{ivk})' \boldsymbol{\psi} + \alpha_i + \lambda_{vk}, \end{aligned} \quad (1)$$

estimated by Poisson pseudo-maximum-likelihood with standard errors clustered on district. Each coefficient vector has twenty-four elements: an observation at position k uses element k , the season that opened its own monsoon year, and element $k + 12$, the season before. This is one regression, not twenty-four — estimating it month by month would lose more than half the observations, since a Poisson drops districts with no conflict in the month being estimated.

What varies, and what does not.—The rainfall on the right-hand side is a seasonal total, so what changes across the twenty-four coefficients is which month’s violence sits on the left; month-to-month movement in the estimates is not structure. The twenty-four positions all fall after the monsoon begins, so the design contains no pre-shock months and offers no placebo lead.

Controls and fixed effects.— $\mathbf{x}_{ivk}$ holds the district’s own rainfall and temperature that month and the month before; $\mathbf{W} \mathbf{z}_{ivk}$ holds the queen-contiguity averages of neighbours’ previous-month conflict and of neighbours’ own weather, rainfall in logs on both sides. The contemporaneous spatial lag is omitted for the reason Table E.7 gives, and

the Poisson carries no own lagged conflict term. The time effect is $\lambda_{vk} = \delta_{vk}$, one parameter for every calendar month in the sample (424 of them): a single month-year effect absorbs anything common to India in that month — an election, a national operation, a change in reporting — where separate month-of-year and year effects absorb only the average season and the average year. It does not absorb the horizon structure, because within a month-year cell every district sits at the same two horizons, so β_k and β_{k+12} are identified from variation across districts in the two vintages’ rainfall. This is the strictest time control the design admits. Two rungs are reported, without the spatial block and with it.

Result.—Figure D.1 plots both rungs over 1990–2025, the window of Tables 2–E.2. Monsoon rainfall is flat or negative through the season itself and runs positive through the following year: summed over months seven to eighteen it gives 5.40 under month-year effects alone and 1.84 with the spatial block, the peak falling in November of the following year (0.58). Non-monsoon rainfall, which is nearly orthogonal to monsoon rainfall once the fixed effects are removed, does nothing over the same window in either rung: -2.89 and -0.91. The delayed response is a monsoon phenomenon rather than a rainfall phenomenon, which is what the agricultural mechanism requires.

What the spatial block carries.—The two rungs differ in size rather than shape. Adding neighbours’ lagged conflict and neighbours’ weather cuts the monsoon sum by roughly two thirds while leaving the profile — negative in season, positive from the turn of the year, peaking in the autumn of $t + 1$ — intact. Part of what the monthly profile picks up without the block is therefore neighbourhood diffusion rather than the district’s own response, and the spatial rung is the one the tables’ column 4 corresponds to and the one the text quotes.




Figure D.1: Monthly horizons of the upstream-rainfall shock, by season

Notes: Each panel plots $\hat{\beta}_h^M$ and $\hat{\beta}_h^N$ from one estimate of equation (1), both seasons carried simultaneously at all twenty-four horizons, against months since June of the shock year. Points are Poisson coefficients with 95 percent cluster-by-district intervals. The panels differ only in the spatial block; both carry district and month-year effects, and 424 month-year cells span 1990–2025. Shaded bands mark the months in which each rainfall measure is accumulated. The monsoon measure is the June–September total of the shock year, so coefficients inside the first band are contemporaneous with the rain rather than responses to it. The second band is the October–December portion of the non-monsoon measure; its remaining five months, January–May of the same calendar year, fall before horizon 0 and lie off the left edge, so most of the non-monsoon treatment is accumulated before the monsoon it is raced against. The third band is the following year’s monsoon, which falls inside the window the annual lag term aggregates (horizons 7–18, between the dashed rules): a lagged annual coefficient cannot separate its own shock’s late response from the next monsoon’s early one, and the monthly profile is what makes that visible. Estimated on 86,496 district-months over 204 districts; the rung without the spatial block runs on 86,496 over 204. Identity is carried by shape as well as colour.

D.2. Agriculture, Prices and Labour

Table D.1: Kharif crop margins: extensive and intensive

	Rice (1)	Cotton (2)	Bajra (3)	Maize (4)	Groundnut (5)	Pigeonpea (6)
<i>Area, extensive: any planting (0/1)</i>						
log Upstream Precip. (<i>t</i>)	-0.060*** (0.014)	-0.090*** (0.027)	-0.058** (0.025)	-0.135*** (0.024)	-0.066** (0.028)	-0.055** (0.026)
log Downstream Precip. (<i>t</i>)	0.072*** (0.016)	-0.011 (0.026)	0.007 (0.022)	0.081*** (0.024)	0.014 (0.031)	0.017 (0.026)
Observations	8,300	8,300	8,300	8,300	8,300	8,300
<i>Area, intensive: log hectares, growing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.035 (0.076)	0.336** (0.131)	-0.049 (0.123)	-0.376*** (0.087)	0.015 (0.130)	-0.256** (0.100)
log Downstream Precip. (<i>t</i>)	0.086 (0.063)	0.028 (0.105)	0.042 (0.114)	0.057 (0.091)	0.433*** (0.117)	0.103 (0.111)
Observations	6,389	3,406	4,944	6,955	4,951	6,648
<i>Production, extensive: any output (0/1)</i>						
log Upstream Precip. (<i>t</i>)	-0.063*** (0.014)	-0.080*** (0.026)	-0.064*** (0.025)	-0.164*** (0.029)	-0.073*** (0.028)	-0.051** (0.025)
log Downstream Precip. (<i>t</i>)	0.070*** (0.016)	-0.020 (0.026)	0.002 (0.022)	0.090*** (0.027)	0.017 (0.032)	0.012 (0.025)
Observations	8,300	8,300	8,300	8,300	8,300	8,300
<i>Production, intensive: log tonnes, producing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.297*** (0.089)	0.665*** (0.183)	0.012 (0.133)	-0.273*** (0.098)	0.092 (0.143)	-0.158 (0.106)
log Downstream Precip. (<i>t</i>)	0.051 (0.070)	-0.003 (0.149)	0.093 (0.153)	0.115 (0.103)	0.402*** (0.111)	0.018 (0.123)
Observations	6,381	3,282	4,900	6,773	4,888	6,563

Notes: Every cell is one regression on the specification of Table D.7's agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Extensive rows are linear probability models on an indicator for any planting (or any output) among district-years where the crop is observed; intensive rows are logs on the district-years that grow (produce) it. Columns are the six largest crops by the season's planted area, grown in at least forty districts. The season is the monsoon's own: rain falls on the standing crop, so the extensive rows measure plantings prevented or destroyed. ***, **, and * denote 1%, 5%, and 10%.

Table D.2: Rabi crop margins: extensive and intensive

	Wheat (1)	Chickpea (2)	Rape & mustard (3)	Jowar (4)	Maize (5)	Barley (6)
<i>Area, extensive: any planting (0/1)</i>						
log Upstream Precip. (<i>t</i>)	0.004 (0.013)	-0.012 (0.021)	-0.012 (0.024)	0.020 (0.017)	0.007 (0.023)	0.002 (0.019)
log Downstream Precip. (<i>t</i>)	0.003 (0.012)	-0.031 (0.025)	0.057*** (0.020)	-0.057*** (0.020)	-0.037* (0.022)	0.007 (0.016)
Observations	8,248	8,248	8,248	8,248	8,248	8,248
<i>Area, intensive: log hectares, growing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.291*** (0.071)	-0.131 (0.103)	-0.037 (0.086)	0.091 (0.206)	-0.038 (0.235)	-0.044 (0.084)
log Downstream Precip. (<i>t</i>)	-0.087 (0.057)	0.105 (0.096)	-0.069 (0.078)	-0.123 (0.200)	0.558*** (0.212)	0.073 (0.095)
Observations	7,759	7,157	7,042	1,173	3,014	4,593
<i>Production, extensive: any output (0/1)</i>						
log Upstream Precip. (<i>t</i>)	0.000 (0.012)	-0.012 (0.022)	-0.001 (0.022)	0.021 (0.017)	0.008 (0.023)	0.011 (0.019)
log Downstream Precip. (<i>t</i>)	0.003 (0.012)	-0.042* (0.024)	0.047** (0.021)	-0.058*** (0.020)	-0.037* (0.022)	-0.004 (0.016)
Observations	8,248	8,248	8,248	8,248	8,248	8,248
<i>Production, intensive: log tonnes, producing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.287*** (0.085)	-0.129 (0.110)	-0.038 (0.094)	0.025 (0.195)	-0.065 (0.239)	0.006 (0.089)
log Downstream Precip. (<i>t</i>)	-0.198*** (0.065)	0.070 (0.104)	-0.112 (0.083)	-0.125 (0.200)	0.710*** (0.219)	0.031 (0.110)
Observations	7,753	7,020	6,994	1,169	3,006	4,492

Notes: Every cell is one regression on the specification of Table D.7's agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Extensive rows are linear probability models on an indicator for any planting (or any output) among district-years where the crop is observed; intensive rows are logs on the district-years that grow (produce) it. Columns are the six largest crops by the season's planted area, grown in at least forty districts. The monsoon precedes rabi sowing, so the contemporaneous term is rain on the soil moisture the winter crop is sown into. Rabi area is concentrated — wheat alone is three-fifths of it — so the smaller columns carry wide intervals. ***, **, and * denote 1%, 5%, and 10%.

Table D.3: Kharif crop prices and the value of the harvest

	Rice (1)	Cotton (2)	Bajra (3)	Maize (4)	Groundnut (5)	Pigeonpea (6)
<i>Farm-harvest price, annual (log)</i>						
log Upstream Precip. (<i>t</i>)	0.004 (0.025)	0.055 (0.038)	-0.042 (0.030)	-0.043* (0.025)	-0.064** (0.030)	0.013 (0.026)
log Downstream Precip. (<i>t</i>)	-0.076*** (0.020)	0.043 (0.045)	0.006 (0.023)	0.018 (0.021)	0.039 (0.029)	0.027 (0.027)
Observations	5,013	1,734	3,429	4,740	3,794	3,912
<i>Revenue: price $\times$ production (log)</i>						
log Upstream Precip. (<i>t</i>)	0.385*** (0.109)	0.168 (0.409)	-0.007 (0.209)	-0.310** (0.127)	-0.027 (0.175)	-0.007 (0.167)
log Downstream Precip. (<i>t</i>)	-0.148* (0.080)	-0.048 (0.215)	0.175 (0.185)	0.268* (0.136)	0.296** (0.133)	0.117 (0.225)
Observations	2,489	887	1,975	2,641	2,133	2,535

Notes: Every cell is one regression on the specification of Table D.7's agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Prices are the DLD farm-harvest series, which is annual — it is not reported by season — so the revenue rows pair the annual price with kharif production: roughly the value of the harvest, and gross revenue rather than profit, since input costs are not observed. Because $\log(pq) = \log p + \log q$, the revenue coefficient equals the sum of the price and production responses on its own sample; it will not equal the sum of the rows printed above it or in Table D.1, which are estimated on the larger samples their own margins allow. The row's contribution is the standard error of that sum, which the covariance of the parts makes unreadable from the separate margins. Both blocks condition on producing district-years with a price quote, so they are intensive margins. ***, **, and * denote 1%, 5%, and 10%.

Table D.4: Kharif crop margins on the conflict-estimation sample

	Rice (1)	Cotton (2)	Bajra (3)	Maize (4)	Groundnut (5)	Pigeonpea (6)
<i>Area, extensive: any planting (0/1)</i>						
log Upstream Precip. (t)	-0.035* (0.019)	-0.037 (0.037)	-0.028 (0.035)	-0.137*** (0.039)	-0.048 (0.043)	-0.018 (0.036)
log Downstream Precip. (t)	0.079*** (0.022)	-0.072** (0.030)	-0.023 (0.030)	0.105*** (0.036)	-0.027 (0.039)	-0.000 (0.036)
Observations	3,841	3,841	3,841	3,841	3,841	3,841
<i>Area, intensive: log hectares, growing district-years</i>						
log Upstream Precip. (t)	-0.008 (0.094)	0.216 (0.194)	-0.130 (0.194)	-0.442*** (0.144)	-0.084 (0.208)	-0.249** (0.124)
log Downstream Precip. (t)	0.282*** (0.080)	-0.044 (0.138)	0.083 (0.147)	0.128 (0.122)	0.431*** (0.163)	0.069 (0.123)
Observations	2,563	1,595	1,901	2,985	1,836	3,072
<i>Production, extensive: any output (0/1)</i>						
log Upstream Precip. (t)	-0.042** (0.018)	-0.028 (0.038)	-0.032 (0.036)	-0.159*** (0.043)	-0.041 (0.046)	-0.023 (0.035)
log Downstream Precip. (t)	0.078*** (0.023)	-0.076** (0.030)	-0.029 (0.031)	0.098** (0.040)	-0.021 (0.040)	-0.003 (0.034)
Observations	3,841	3,841	3,841	3,841	3,841	3,841
<i>Production, intensive: log tonnes, producing district-years</i>						
log Upstream Precip. (t)	0.392*** (0.122)	0.860*** (0.280)	0.018 (0.197)	-0.174 (0.177)	0.161 (0.207)	-0.083 (0.153)
log Downstream Precip. (t)	0.282*** (0.090)	-0.015 (0.214)	0.106 (0.188)	0.208 (0.140)	0.458*** (0.155)	0.054 (0.142)
Observations	2,559	1,544	1,890	2,883	1,803	3,048

Notes: Every cell is one regression on the specification of Table D.7's agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Extensive rows are linear probability models on an indicator for any planting (or any output) among district-years where the crop is observed; intensive rows are logs on the district-years that grow (produce) it. Columns are the six largest crops by the season's planted area, grown in at least forty districts. The sample is restricted to the districts the conflict Poisson retains — at least one located UCDP event over 1990–2025; columns are kept identical to the full-sample table for comparability. ***, **, and * denote 1%, 5%, and 10%.

Table D.5: Rabi crop margins on the conflict-estimation sample

	Wheat (1)	Chickpea (2)	Rape & mustard (3)	Jowar (4)	Maize (5)	Barley (6)
<i>Area, extensive: any planting (0/1)</i>						
log Upstream Precip. (<i>t</i>)	-0.040** (0.019)	-0.037 (0.033)	-0.047* (0.028)	-0.007 (0.022)	-0.030 (0.038)	-0.011 (0.033)
log Downstream Precip. (<i>t</i>)	0.026** (0.013)	-0.060* (0.031)	0.034 (0.022)	-0.026 (0.025)	-0.027 (0.026)	0.006 (0.023)
Observations	3,903	3,903	3,903	3,903	3,903	3,903
<i>Area, intensive: log hectares, growing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.293** (0.122)	-0.245 (0.153)	-0.345*** (0.113)	0.151 (0.581)	-0.291 (0.362)	0.072 (0.133)
log Downstream Precip. (<i>t</i>)	-0.169** (0.075)	0.130 (0.107)	-0.050 (0.079)	0.153 (0.458)	1.160*** (0.355)	-0.008 (0.112)
Observations	3,785	3,373	3,564	299	1,451	2,107
<i>Production, extensive: any output (0/1)</i>						
log Upstream Precip. (<i>t</i>)	-0.040** (0.019)	-0.038 (0.033)	-0.053** (0.027)	-0.007 (0.022)	-0.028 (0.038)	-0.001 (0.032)
log Downstream Precip. (<i>t</i>)	0.026** (0.013)	-0.066** (0.032)	0.029 (0.022)	-0.026 (0.025)	-0.026 (0.026)	-0.008 (0.024)
Observations	3,903	3,903	3,903	3,903	3,903	3,903
<i>Production, intensive: log tonnes, producing district-years</i>						
log Upstream Precip. (<i>t</i>)	0.206 (0.140)	-0.282* (0.167)	-0.416*** (0.112)	0.104 (0.509)	-0.216 (0.359)	0.091 (0.141)
log Downstream Precip. (<i>t</i>)	-0.261*** (0.085)	0.111 (0.112)	-0.134 (0.085)	0.331 (0.434)	1.252*** (0.361)	-0.031 (0.132)
Observations	3,785	3,330	3,557	298	1,448	2,073

Notes: Every cell is one regression on the specification of Table D.7's agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Extensive rows are linear probability models on an indicator for any planting (or any output) among district-years where the crop is observed; intensive rows are logs on the district-years that grow (produce) it. Columns are the six largest crops by the season's planted area, grown in at least forty districts. The sample is restricted to the districts the conflict Poisson retains — at least one located UCDP event over 1990–2025; columns are kept identical to the full-sample table for comparability. Rabi jowar falls to 33 growing districts here, below the forty-district floor the full-sample selection imposes; the column is retained for comparability and should be read accordingly. ***, **, and * denote 1%, 5%, and 10%.

Table D.6: Kharif crop prices and the value of the harvest on the conflict-estimation sample

	Rice (1)	Cotton (2)	Bajra (3)	Maize (4)	Groundnut (5)	Pigeonpea (6)
<i>Farm-harvest price, annual (log)</i>						
log Upstream Precip. (<i>t</i>)	0.026 (0.036)	0.046 (0.055)	-0.033 (0.049)	-0.043 (0.034)	-0.045 (0.039)	0.106*** (0.034)
log Downstream Precip. (<i>t</i>)	-0.080*** (0.023)	0.110* (0.064)	-0.006 (0.038)	-0.004 (0.027)	0.047 (0.041)	0.026 (0.043)
Observations	2,794	896	1,414	2,377	1,645	1,871
<i>Revenue: price × production (log)</i>						
log Upstream Precip. (<i>t</i>)	0.349** (0.139)	0.222 (0.552)	-0.258 (0.282)	-0.175 (0.202)	-0.060 (0.326)	0.113 (0.229)
log Downstream Precip. (<i>t</i>)	0.002 (0.095)	-0.132 (0.316)	0.331 (0.274)	0.273* (0.160)	0.500*** (0.181)	0.487* (0.289)
Observations	1,043	473	689	1,083	743	1,172

Notes: Every cell is one regression on the specification of Table D.7’s agricultural rows — upstream, local, neighbour and downstream monsoon rainfall, each contemporaneous and lagged, temperature controls, district and year fixed effects, standard errors clustered on district; the contemporaneous upstream and downstream terms are shown. Prices are the DLD farm-harvest series, which is annual — it is not reported by season — so the revenue rows pair the annual price with kharif production: roughly the value of the harvest, and gross revenue rather than profit, since input costs are not observed. Because $\log(pq) = \log p + \log q$, the revenue coefficient equals the sum of the price and production responses on its own sample; it will not equal the sum of the rows printed above it or in Table D.1, which are estimated on the larger samples their own margins allow. The row’s contribution is the standard error of that sum, which the covariance of the parts makes unreadable from the separate margins. Both blocks condition on producing district-years with a price quote, so they are intensive margins. ***, **, and * denote 1%, 5%, and 10%. The sample is restricted to the districts the conflict Poisson retains — at least one located UCDP event over 1990–2025; columns are kept identical to the full-sample table for comparability.

Table D.7: Four sources of water in one regression

	Each outcome's own sample				Conflict-estimation districts			
	Upstream (1)	Local (2)	Neighbour (3)	Downstream (4)	Upstream (5)	Local (6)	Neighbour (7)	Downstream (8)
<i>A. Conflict</i>								
Conflict events (Poisson) ($t - 1$)	1.144*** (0.324)	0.622 (0.910)	-1.413 (1.260)	-0.977*** (0.282)	1.144*** (0.324)	0.622 (0.910)	-1.413 (1.260)	-0.977*** (0.282)
<i>B. District agriculture, prices and wages</i>								
Kharif planted area (t)	-0.102 (0.071)	0.033 (0.168)	0.181 (0.206)	0.217*** (0.056)	-0.343*** (0.128)	0.063 (0.268)	0.372 (0.343)	0.337*** (0.087)
Fixed-weight quantity index ($t - 1$)	0.159* (0.081)	-0.144 (0.194)	0.147 (0.226)	0.118** (0.054)	0.006 (0.143)	-0.035 (0.328)	0.022 (0.385)	0.174** (0.085)
Implicit price index ($t - 1$)	-0.057** (0.022)	0.036 (0.051)	-0.047 (0.056)	0.029 (0.021)	-0.061** (0.026)	0.061 (0.076)	-0.068 (0.085)	-0.009 (0.029)
Fixed-weight price index (t)	-0.012 (0.021)	0.019 (0.045)	-0.042 (0.058)	0.001 (0.016)	-0.007 (0.031)	0.050 (0.071)	-0.050 (0.094)	-0.001 (0.021)
Upstream districts' output (t)	0.269*** (0.049)	-0.010 (0.120)	0.099 (0.145)	-0.011 (0.037)	0.189*** (0.067)	0.009 (0.156)	0.172 (0.184)	-0.094** (0.046)
Real agricultural wage ($t - 1$)	-0.127*** (0.041)	-0.004 (0.077)	0.026 (0.100)	0.094*** (0.032)	-0.108** (0.054)	-0.045 (0.095)	0.135 (0.127)	0.087** (0.038)
<i>C. IHDS households</i>								
Farm income (asinh) (t)	0.305 (1.529)	-3.365* (1.715)	2.620 (2.068)	3.242*** (1.152)	4.144 (2.645)	-4.808** (1.925)	0.265 (2.044)	2.903** (1.312)
Non-farm income (asinh) ($t - 1$)	-3.424** (1.430)	1.472 (1.786)	1.183 (2.539)	-0.167 (0.827)	-5.905*** (2.201)	4.087 (2.730)	2.613 (2.921)	-0.472 (1.058)
Receives any remittance ($t - 1$)	0.075 (0.064)	0.038 (0.118)	-0.248* (0.131)	-0.044 (0.044)	0.013 (0.104)	0.139 (0.155)	-0.325* (0.187)	-0.015 (0.070)
Consumption (log) (t)	-0.650*** (0.177)	0.559** (0.217)	-0.077 (0.247)	0.110 (0.111)	0.213 (0.396)	-0.316 (0.339)	0.202 (0.280)	-0.289* (0.163)
Districts, conflict row	177				177			
District & Year (or Wave) FE	✓				✓			

Notes: Every row is one regression carrying all four rainfall measures at once, each entered contemporaneously and lagged, with upstream and local temperatures and their lags, district and year (or survey-wave) fixed effects, and standard errors clustered on district. Each row reports one horizon, marked against its label: ($t - 1$) for the lagged term and (t) for the contemporaneous one. Crop quantities and the prices they set are read contemporaneously, since a crop is harvested in the year the rain falls; conflict, wages and the slower household margins are read at $t - 1$. Upstream is the river-distance-weighted catchment draining into the district excluding the district and its neighbours; downstream mirrors it exactly, on the catchment the district drains into; neighbour is the row-normalised queen-contiguity average of local rainfall. Columns 1–4 estimate each outcome on whatever districts it supports; columns 5–8 restrict every outcome to the 177 districts the conflict Poisson retains, so the comparison across outcome families holds geography fixed. The crop aggregates respond at the extensive margin — plantings pushed to zero in district-years where the crop is marginal — which Tables D.1 and D.2 decompose by season; the response among surviving plantings is flat to positive. Local and neighbour rainfall correlate 0.964 after fixed effects and carry variance inflation factors of 14.0 and 16.4, so neither is separately identified; upstream and downstream carry 2.50 and 2.26.




Figure D.2: The dose–response in upstream and downstream rainfall anomalies

Notes: One fixed-effects Poisson regression carrying both sources as five-bin functions of their own standardised monsoon anomaly, with the middle bin (-0.5 to 0.5 standard deviations, shaded) omitted, so every coefficient reads against a normal year. Bins rather than a polynomial: a quadratic imposes a single turning point and would conceal a one-sided response, which is the shape in question. Both terms are lagged, as in the headline. Controls are the local rainfall anomaly and the upstream, downstream and local temperature anomalies, each current and lagged, with district and year effects; standard errors are clustered on district and bars are 95 percent intervals. The control set is the anomaly counterpart of Table D.7 rather than that table’s right-hand side exactly — it carries no neighbour block — so the two are estimated on the same sample and the same sources but are not nested. 6,372 district-years over 177 districts. Both panels share a vertical scale, since the comparison the figure exists to make is one of shape *and* magnitude.

Table D.8: Conflict and the mechanism outcomes in rainfall anomalies

	Conflict events (1)	Planted area (2)	Crops planted (3)	Quantity index (4)	Implicit price (5)	Real wage (6)	Nominal wage (7)
<i>Upstream anomaly, contemporaneous (t)</i>							
< -1.5 SD	0.231 (0.197)	-0.084 (0.056)	0.126** (0.058)	-0.163** (0.064)	0.016 (0.015)	-0.042 (0.027)	-0.073*** (0.023)
-1.5 to -0.5	-0.118 (0.118)	0.097*** (0.026)	0.094*** (0.029)	0.061** (0.029)	0.017** (0.007)	-0.003 (0.011)	-0.012 (0.009)
0.5 to 1.5	0.059 (0.138)	0.018 (0.031)	-0.060** (0.027)	0.041 (0.033)	0.004 (0.008)	-0.033** (0.015)	-0.008 (0.014)
> 1.5 SD	0.217 (0.189)	0.045 (0.045)	-0.113*** (0.039)	-0.006 (0.046)	0.033** (0.013)	-0.058** (0.027)	-0.035 (0.021)
<i>Upstream anomaly, lagged ($t - 1$)</i>							
< -1.5 SD	-0.289* (0.162)	-0.012 (0.071)	-0.011 (0.053)	0.013 (0.073)	0.023 (0.014)	0.027 (0.028)	-0.020 (0.022)
-1.5 to -0.5	-0.048 (0.105)	0.032 (0.024)	0.036 (0.026)	0.004 (0.028)	0.017** (0.007)	0.002 (0.012)	-0.009 (0.010)
0.5 to 1.5	0.368*** (0.095)	0.069** (0.029)	-0.018 (0.025)	0.050 (0.032)	-0.017** (0.008)	-0.025 (0.016)	-0.010 (0.013)
> 1.5 SD	0.229 (0.154)	0.109** (0.047)	-0.115** (0.049)	0.063 (0.051)	0.018 (0.012)	-0.056** (0.027)	-0.049** (0.023)
<i>Downstream anomaly, contemporaneous (t)</i>							
< -1.5 SD	-0.296 (0.201)	-0.158*** (0.042)	-0.151** (0.065)	-0.123*** (0.044)	0.020 (0.015)	-0.024 (0.024)	0.017 (0.020)
-1.5 to -0.5	0.307*** (0.106)	-0.112*** (0.022)	-0.079*** (0.028)	-0.128*** (0.027)	0.031*** (0.007)	-0.018 (0.014)	0.008 (0.012)
0.5 to 1.5	-0.064 (0.111)	0.037* (0.022)	0.044** (0.022)	0.062*** (0.023)	0.002 (0.007)	0.002 (0.013)	-0.006 (0.010)
> 1.5 SD	-0.221* (0.127)	0.192*** (0.053)	0.136*** (0.041)	0.195*** (0.053)	0.030** (0.012)	0.008 (0.023)	-0.004 (0.018)
<i>Downstream anomaly, lagged ($t - 1$)</i>							
< -1.5 SD	-0.043 (0.216)	0.033 (0.041)	-0.085 (0.055)	0.048 (0.048)	-0.019 (0.014)	-0.030 (0.025)	-0.006 (0.022)
-1.5 to -0.5	0.199* (0.112)	-0.051** (0.023)	-0.029 (0.026)	-0.038 (0.027)	0.016** (0.007)	-0.021 (0.015)	-0.006 (0.013)
0.5 to 1.5	-0.291** (0.141)	0.031 (0.020)	-0.048** (0.024)	0.045* (0.024)	0.022*** (0.008)	0.000 (0.014)	0.006 (0.011)
> 1.5 SD	-0.475*** (0.169)	0.177*** (0.052)	0.070 (0.052)	0.147** (0.057)	0.022 (0.013)	0.028 (0.025)	0.017 (0.021)
Estimation	Poisson	OLS	OLS	OLS	OLS	OLS	OLS
Observations	6,372	8,079	8,300	8,079	5,806	3,723	3,723
Districts	177	341	342	341	263	276	276

Notes: Every column is one regression in which upstream, downstream and local monsoon rainfall enter as unrestricted bin functions of each basin-set's own standardised anomaly — the deviation from the district's long-run mean divided by its own standard deviation — at both horizons jointly, with contemporaneous and lagged temperature anomalies as controls and district and year fixed effects; standard errors are clustered on district. The normal bin (−0.5 to 0.5 SD) is omitted, so every coefficient reads against a normal year; local rainfall's bins are controls and are not shown. Mechanism outcomes use the six-crop constructions of Table D.7 on each outcome's own sample, 1997–2022, by OLS; conflict is the fixed-effects Poisson on located events over 1990–2025, the paper's estimation sample. Crops planted counts how many of the six kharif crops have positive area, on district-years observing all six. ***, **, and * denote 1%, 5%, and 10%.




Figure D.3: Conflict and agriculture across the rainfall anomaly

Notes: Points are coefficients on bins of the lagged upstream rainfall anomaly, relative to the omitted -0.5 to 0.5 bin, with 95 percent cluster-by-district intervals. The anomaly is a within-district z -score computed on each district's own annual history. The conflict panel is Poisson on located events over 1990–2025 under the specification of Table 2, column 4 — the same sample of 8,280 district-years and 230 districts as the headline. The agricultural panels are OLS on log outcomes over the years those series cover. Bins rather than a quadratic, because a quadratic imposes one turning point and lets the tails drive the curvature through leverage.

Table D.9: The wage response is a multi-year adjustment

	Real wage (1)	Nominal wage (2)
log Upstream Precip. ($t - 1$)	-0.065 (0.042)	-0.050* (0.031)
log Upstream Precip. ($t - 2$)	-0.046** (0.023)	-0.012 (0.023)
log Upstream Precip. ($t - 3$)	-0.035 (0.028)	-0.037 (0.024)
log Upstream Precip. ($t - 4$)	0.034 (0.030)	-0.010 (0.025)
log Upstream Precip. ($t - 5$)	-0.012 (0.029)	-0.020 (0.024)
Cumulative, $\sum_{k=1}^5$	-0.123 (0.103)	-0.129 (0.090)
Observations	3,872	3,872
Districts	276	276

Notes: Five annual lags of log upstream monsoon rainfall entered together, with local, neighbour and downstream rainfall and the temperature controls as in Table 6, district and year fixed effects, and standard errors clustered on district. The cumulative row is the sum of the five lags with the standard error of that sum. Table 6 reports the single lagged term, which attributes to one year a response that arrives over three: the first three lags each carry part of it and the fourth and fifth carry none. The response is therefore weaker than the single-lag coefficient implies, and it is not drift alone: individual years carry it, which two smoothly co-moving series would not produce. The distributional test the wage series invites — female field labour is disproportionately landless — cannot be run: the female series collapses after 2004 and concentrates in five states, so a female-minus-male gap reports on the sample, not on incidence. ***, **, and * denote 1%, 5%, and 10%.

Table D.10: Labour supply and earnings after an upstream shock: NSS agricultural margins, by fieldwork quarter

	Any ag. wage work (1)	Days, ag. wage work (2)	Ag. wage earnings (3)	All wage earnings (4)
<i>All households</i>				<i>n</i> = 547,606
log Upstream Precip. (<i>t</i>)	+0.0160** (0.0076)	+0.1012** (0.0456)	+0.1143** (0.0476)	-0.0591 (0.0819)
log Upstream Precip. (<i>t</i> − 1)	+0.0061 (0.0091)	+0.0182 (0.0545)	+0.0828 (0.0562)	+0.0842 (0.1098)
<i>Interviewed Jul–Sep, during the monsoon</i>				<i>n</i> = 137,591
log Upstream Precip. (<i>t</i>)	+0.0133 (0.0136)	+0.1026 (0.0772)	+0.0996 (0.0870)	-0.1628 (0.1367)
log Upstream Precip. (<i>t</i> − 1)	-0.0109 (0.0156)	-0.0885 (0.0910)	-0.0424 (0.0978)	-0.1817 (0.1692)
<i>Interviewed Oct–Dec, kharif harvest (+2 months)</i>				<i>n</i> = 137,619
log Upstream Precip. (<i>t</i>)	+0.0153 (0.0130)	+0.0818 (0.0794)	+0.1040 (0.0818)	+0.0198 (0.1361)
log Upstream Precip. (<i>t</i> − 1)	+0.0025 (0.0162)	-0.0344 (0.0965)	+0.0683 (0.1021)	+0.3293* (0.1701)
<i>Interviewed Jan–Mar (+5 months)</i>				<i>n</i> = 135,937
log Upstream Precip. (<i>t</i>)	+0.0070 (0.0121)	+0.0526 (0.0726)	+0.0572 (0.0798)	-0.2774* (0.1452)
log Upstream Precip. (<i>t</i> − 1)	+0.0029 (0.0154)	+0.0163 (0.0942)	+0.0395 (0.0993)	-0.2410 (0.1784)
<i>Interviewed Apr–Jun, following year (+8 months)</i>				<i>n</i> = 136,451
log Upstream Precip. (<i>t</i>)	+0.0283*** (0.0103)	+0.1671*** (0.0642)	+0.1969*** (0.0664)	+0.1559 (0.1453)
log Upstream Precip. (<i>t</i> − 1)	+0.0295** (0.0133)	+0.1776** (0.0811)	+0.2655*** (0.0864)	+0.4037** (0.1740)

Notes: NSS Employment–Unemployment rounds 61, 66 and 68 (2004–05, 2009–10, 2011–12), working-age individuals, on the four sources of water — upstream, local, neighbouring and downstream rainfall — with temperatures, each at the contemporaneous and lagged monsoon, district and round fixed effects, individual age, sex and education, and standard errors clustered on district. Each round follows exactly one monsoon — 2004, 2009 and 2011 — which is the contemporaneous term. THE REFERENCE PERIOD IS ONE WEEK: days run 0–7. Earnings are wage and salary receipts for that same week in asinh, which keeps the structural zeros — most people do no wage work of a given kind in a given week, and the margin that moves is whether they do any — but prices the jump from zero in rupees, so the coefficients read as sign and significance rather than as elasticities. NSS allocates households to fieldwork quarters BY DESIGN, so the panels compare households drawn to be alike rather than households that selected their timing. The quarters sit at different distances from the SAME monsoon, and the panel headings give that distance: July–September is surveyed during it and April–June eight months after it ended, in the following calendar year — every panel is post-treatment. Estimates without the demographic controls are close to these and are in the replication output. ***, **, and * denote 1%, 5%, and 10%. A coefficient of 0.15 on days is about eight days a year if the response were constant across the year, which the panels show it is not. The response is absent in the first two quarters and concentrated in the last, which is the calendar Section VI reports for conflict. Table D.11 reports the non-farm margins on the same specification.

Table D.11: The non-farm margin after an upstream shock: NSS, by fieldwork quarter

	Casual non-farm			Self-employment		Regular salaried
	Any work (1)	Days (2)	Earnings (3)	Any (4)	Days (5)	Earnings (6)
<i>All households</i>						<i>n</i> = 547,606
log Upstream Precip. (<i>t</i>)	-0.0155 (0.0100)	-0.0884 (0.0663)	-0.1172* (0.0670)	+0.0215 (0.0226)	+0.2485* (0.1472)	-0.0600 (0.0706)
log Upstream Precip. (<i>t</i> − 1)	-0.0298*** (0.0115)	-0.1855** (0.0736)	-0.2033** (0.0803)	+0.0330 (0.0239)	+0.2474 (0.1549)	+0.1959** (0.0783)
<i>Interviewed Jul–Sep, during the monsoon</i>						<i>n</i> = 137,591
log Upstream Precip. (<i>t</i>)	-0.0294* (0.0162)	-0.1320 (0.1024)	-0.2241** (0.1013)	+0.0118 (0.0318)	+0.1660 (0.2081)	-0.0561 (0.1059)
log Upstream Precip. (<i>t</i> − 1)	-0.0142 (0.0140)	-0.0844 (0.0874)	-0.0812 (0.0973)	+0.0642* (0.0375)	+0.5585** (0.2338)	-0.0697 (0.1222)
<i>Interviewed Oct–Dec, kharif harvest (+2 months)</i>						<i>n</i> = 137,619
log Upstream Precip. (<i>t</i>)	-0.0127 (0.0120)	-0.0900 (0.0755)	-0.0875 (0.0845)	+0.0169 (0.0312)	+0.2248 (0.2052)	+0.0058 (0.1012)
log Upstream Precip. (<i>t</i> − 1)	-0.0337** (0.0154)	-0.1813* (0.0946)	-0.2085* (0.1069)	+0.0067 (0.0320)	+0.0527 (0.2095)	+0.4547*** (0.1374)
<i>Interviewed Jan–Mar (+5 months)</i>						<i>n</i> = 135,937
log Upstream Precip. (<i>t</i>)	-0.0254* (0.0150)	-0.1556 (0.0982)	-0.1920* (0.1046)	+0.0603** (0.0288)	+0.4901*** (0.1737)	-0.1438 (0.1138)
log Upstream Precip. (<i>t</i> − 1)	-0.0455*** (0.0170)	-0.2740** (0.1097)	-0.3339*** (0.1188)	+0.0671** (0.0340)	+0.4614** (0.2061)	+0.0489 (0.1218)
<i>Interviewed Apr–Jun, following year (+8 months)</i>						<i>n</i> = 136,451
log Upstream Precip. (<i>t</i>)	+0.0050 (0.0139)	+0.0196 (0.0900)	+0.0322 (0.0993)	+0.0010 (0.0322)	+0.1448 (0.2065)	-0.0717 (0.1240)
log Upstream Precip. (<i>t</i> − 1)	-0.0265 (0.0203)	-0.2026 (0.1259)	-0.1942 (0.1422)	-0.0025 (0.0363)	-0.0634 (0.2489)	+0.3281** (0.1482)

Notes: NSS Employment–Unemployment rounds 61, 66 and 68 (2004–05, 2009–10, 2011–12), working-age individuals, on the four sources of water — upstream, local, neighbouring and downstream rainfall — with temperatures, each at the contemporaneous and lagged monsoon, district and round fixed effects, individual age, sex and education, and standard errors clustered on district. Each round follows exactly one monsoon — 2004, 2009 and 2011 — which is the contemporaneous term. THE REFERENCE PERIOD IS ONE WEEK: days run 0–7. Earnings are wage and salary receipts for that same week in asinh, which keeps the structural zeros — most people do no wage work of a given kind in a given week, and the margin that moves is whether they do any — but prices the jump from zero in rupees, so the coefficients read as sign and significance rather than as elasticities. NSS allocates households to fieldwork quarters BY DESIGN, so the panels compare households drawn to be alike rather than households that selected their timing. The quarters sit at different distances from the SAME monsoon, and the panel headings give that distance: July–September is surveyed during it and April–June eight months after it ended, in the following calendar year — every panel is post-treatment. Estimates without the demographic controls are close to these and are in the replication output. ***, **, and * denote 1%, 5%, and 10%. This table is the non-farm counterpart to Table D.10, which reports the agricultural margins on the same specification and sample. Read together they are the two sides of the household’s response: agricultural wage work expands while the casual non-farm option contracts, so the shock is not absorbed by reallocation into non-farm work. Self-employment days are unconditional and include own-cultivation, whose EARNINGS the Employment–Unemployment schedule does not collect — the total in Table D.10 is therefore total WAGE earnings, not total income.

Table D.12: The four sources on the Economic Census, three rounds

	Upstream (1)	Local (2)	Neighbour (3)	Downstream (4)
Establishments	0.256* (0.134)	-0.459 (0.344)	-0.070 (0.463)	0.124 (0.174)
Non-farm employment	0.222 (0.136)	-0.115 (0.357)	-0.057 (0.496)	0.168 (0.186)
Manufacturing employment	-0.140 (0.210)	-0.550 (0.628)	0.897 (0.844)	-0.577** (0.265)
Services employment	0.305** (0.130)	0.011 (0.303)	-0.395 (0.435)	0.346** (0.170)
Observations			1,080	
Districts			360	
District & Round FE			✓	

Notes: The specification of Table D.7 on the Economic Census, whose rounds give three within-district differences against the household survey's one. Each round's treatment is the mean over the years between it and the previous round, and the lagged term is the mean over the window before that. The 1990 round is excluded: its lag window begins in 1981, before the rainfall panel, so every 1990 observation would lose its lagged treatment. The panel is balanced across the three remaining rounds. Upstream rainfall raises establishment counts and non-farm employment, which is the entry margin Section VI reads as diversification out of damaged agriculture. Downstream rainfall does nothing here, so the upstream-downstream mirror documented in Table D.7 is a feature of the agricultural margin rather than of non-farm entry.

Table D.13: Nightlights and the four sources of water, VIIRS

	Mean light (1)	Total light (2)
log Upstream Precip. (t)	-0.043 (0.050)	-0.041 (0.050)
log Upstream Precip. ($t - 1$)	0.072 (0.056)	0.070 (0.057)
log Local Precip. (t)	0.067 (0.070)	0.059 (0.069)
log Local Precip. ($t - 1$)	-0.023 (0.066)	-0.029 (0.066)
log Downstream Precip. (t)	0.240** (0.096)	0.241** (0.096)
log Downstream Precip. ($t - 1$)	0.251** (0.098)	0.252** (0.099)
Observations	3,240	3,240
Number of Districts	360	360
Estimation	OLS	OLS
District & Year FE	✓	✓

Notes: District-year panel, 2013 – 2021 — the VIIRS record. The outcome is the log of district mean radiance in column 1 and of total radiance in column 2. Each column enters all four sources of water under the specification of Table 2: upstream, local and downstream monsoon rainfall, current and lagged, with their temperature controls and the spatial block is deliberately ABSENT here. Neighbours' local rainfall correlates 0.96 with the district's own after fixed effects, so with both in the equation neither is identified; dropping the neighbour terms costs that control and makes the local coefficient readable, which is the point of the column. VIIRS is reported alone because it and DMSP disagree on the sign of the downstream term, and VIIRS is the calibrated instrument: it carries on-board calibration, does not saturate, and resolves roughly 500 metres, whereas DMSP is a six-bit sensor top-coded at 63 whose inter-satellite drift moves precisely the within-district across-year variation this design uses. The cost is the window: VIIRS begins in 2012, so these columns cannot speak to the earlier decades the rest of the paper covers. Conley spatial-HAC standard errors (150 km cutoff) in parentheses. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.

Table D.14: Is the downstream effect a dam-release effect?

<i>A. Downstream rainfall interacted with dam presence</i>			
	Downstream rain, no dam (1)	× dam upstream (2)	Implied effect with a dam (3)
Conflict events (Poisson)	-1.745*** (0.399)	1.410** (0.594)	-0.335 —
Kharif cereal area	0.230* (0.118)	0.113 (0.132)	0.343 —
Kharif cereal production	0.128 (0.128)	0.253* (0.141)	0.381 —
Real agricultural wage	0.073* (0.040)	0.030 (0.037)	0.103 —
<i>B. The same, with an agricultural interaction alongside</i>			
	Downstream rain (1)	× dam upstream (2)	× agriculture (3)
Irrigated share, 2000	-1.763*** (0.427)	1.424** (0.594)	-0.242 (0.315)
Cropped-area share, 1990s	-1.721*** (0.405)	1.453** (0.583)	-0.256 (0.236)
Districts	177, of which 68% have a dam		
District & Year FE	✓		

Notes: Panel A adds an interaction between downstream rainfall (both horizons) and an indicator for any irrigation dam completed by 1999 in the district's upstream catchment, to the four-source specification of Table D.7. Main effects are retained. The moderator is fixed before the estimation window, so it cannot respond to the rainfall shocks being studied, and it is time-invariant, so district effects absorb its level. A dam-release account predicts column 2 to carry the effect — the mechanism needs a reservoir whose release rule can relax when its command area gets its own rain. It runs the other way: the downstream effect is largest where there is no dam and is attenuated where there is one. Dam placement is not random, so column 2 cannot be read as the causal effect of a dam; but the sign is the opposite of the one the hypothesis requires. Panel B asks whether the dam interaction is really an agricultural one: irrigation dams are built where there is irrigable land, so an interaction with dam presence could pick up a larger downstream effect in agricultural districts. Both interactions are entered together, with the agricultural moderator standardised so the main term is the effect at its mean. The dam interaction is unchanged and the agricultural interactions are individually insignificant, so the two are not proxies for one another — the moderators correlate 0.22 and 0.10. Their signs differ: the downstream effect is if anything LARGER in more agricultural districts, which the income reading predicts, and smaller where there is a dam, which the release-rule reading does not.

D.3. Public Employment

Table D.15: The employment guarantee and local rainfall: replication and measure swaps

	Rain, pre-NREGA	$\times$ NREGA	Rain, post-NREGA	Districts
His events, his rainfall	−1.402*** (0.276)	+1.298*** (0.329)	−0.103 (0.158)	239
His events, CHIRPS rainfall	−1.328*** (0.248)	+1.539*** (0.310)	+0.211 (0.161)	239
UCDP events, his rainfall	−0.632 (0.413)	+0.845* (0.493)	+0.213 (0.251)	211
UCDP events, CHIRPS rainfall	−1.477*** (0.361)	+1.698*** (0.425)	+0.220 (0.234)	211

Notes: [Fetzer \(2020\)](#)’s reduced form — local monsoon rainfall interacted with district NREGA availability, his sample and fixed effects — re-estimated here, then with his TRMM rainfall replaced by this paper’s CHIRPS measure, his news-based event counts replaced by UCDP, and both. The moderation replicates in each: rainfall’s pacifying effect before the programme, absent after. The post-programme interval bounds what can be claimed — it rules out a post effect half the pre-programme size but not a quarter of it, so the programme roughly halves the local rainfall–conflict relationship rather than eliminating it. Standard errors clustered on district.

Table D.16: NREGA and the two rainfall shocks under tightening identification

	Local rainfall		Upstream rainfall	
	$\times$ NREGA	Rain, post	$\times$ NREGA	Rain, post
<i>A. Post-programme interaction, 1990–2025</i>				
District + year	+0.511(0.394)	+0.178(0.268)	+1.552(0.388) ***	+1.290(0.350) ***
+ phase $\times$ year	+0.486(0.462)	+0.013(0.283)	+1.828(0.431) ***	+1.239(0.326) ***
+ state $\times$ year	+0.436(0.240)*	+0.165(0.324)	+0.294(0.353)	+0.719(0.378)*
+ state $\times$ phase $\times$ year	+0.299(0.343)	+0.093(0.315)	+0.405(0.290)	+0.644(0.450)
<i>B. Staggered roll-out: NREGA $\times$ rain within region-year</i>				
Full window	−0.028(0.044)		−0.077(0.073)	
2004–2010	−0.057(0.034)*		−0.081(0.057)	
2006–2008, staggering live	−0.063(0.033)*		−0.088(0.036) **	

Notes: Panel A interacts each rainfall source (lagged, logged) with an indicator for NREGA being active in the district, on the fixed-effects Poisson of Table 2, and tightens the fixed effects one step at a time; shown are the interaction and the implied post-programme rain coefficient. Under district and year effects the upstream interaction is large — but “after NREGA” is also simply “later.” Panel B removes “later” entirely: region $\times$ year effects with phase-specific rain slopes, so the interaction is identified only from roll-out timing within the same region and year. What survives is near zero for both sources — the programme neither amplifies nor absorbs either shock on this margin. The 2006–2008 window is where phases 1–3 differ within a year; its sample is small and the coefficients are correspondingly fragile. Standard errors clustered on district. The elasticity pair the text quotes (1.344 to 1.252) is the upstream response before and after within region-year, from the same specification as Panel B.

Appendix E. Additional Figures and Tables

Exhibits are grouped by the main-text section they support; each carries its own notes.

E.1. Descriptive Figures




Figure E.1: UCDP conflict events in India, 1989–2025

Notes: Panel A plots all events against the subset UCDP locates to a district or better, the outcome used throughout the paper; the two are drawn on one count axis, and the gap between them is the share of events UCDP could place no more finely than a state. Panel B splits all events by conflict system, assigned from the UCDP conflict name (Maoist, Kashmir, Northeast, and a residual holding India–Pakistan interstate events, communal non-state violence, and unaffiliated one-sided violence). Panel C splits them by UCDP’s type of violence. Series are distinguished by line style as well as colour, so the figure reads in greyscale.




Figure E.2: Composition of UCDP events in India, 1989–2025

Notes: Bars give each category's share of all events over the full record; the tick on each bar gives its share of the located subset the estimation uses. Bars and ticks coincide to within 1.5 percentage points everywhere, which is the point of plotting them together: the precision restriction is not selective on conflict system or on type of violence. That neutrality is about composition only: the same restriction is highly selective across *districts* — state-located events pile onto one representative point per state — which is what drives the Kashmir result. Categories are as in Figure E.1.




Figure E.3: Mix of violence types within each conflict system, located events, 1990–2025

Notes: Figure E.2 gives each category's share of all events; this figure gives the cross-tabulation, which is what Tables E.4 and 3 condition on. Bars are each system's located events divided among UCDP's three types of violence and sum to 100 percent; the count beside each bar is the events it rests on. Two of the rows are definitional rather than empirical: communal violence is non-state by construction, which is why it contributes no column to Table 3, and the unaffiliated residual is one-sided by construction. Among the three systems the paper estimates, Maoist violence is the most evenly divided (56.3 percent state-based against 43.5 percent one-sided), which is why its pooled count moves least: the two components respond in opposite directions and largely cancel. Kashmir is 80.0 percent state-based and the Northeast 65.4 percent. Restricting to located events changes no cell of this table by more than 2.4 percentage points, so the composition here is not an artifact of the precision rule. Every segment carries its value, so the figure reads in greyscale.




Figure E.4: Monthly composition of conflict and rainfall, 1990–2025

Notes: Each series is normalised to sum to 100 percent across the twelve calendar months, so the panels show *when* each series occurs rather than how large it is. Panels A and B use located UCDP events; violence types are UCDP's state-based, non-state and one-sided categories, and conflict systems are the paper's Maoist, Kashmir and Northeast groupings with everything else pooled into Other. Panel C rebuilds upstream rainfall from the basin-month series through the same crosswalk and the same river-distance weights that the annual treatment uses, so the seasonal shape shown is the treatment's own; district monthly shares are computed first and then averaged across districts, so that a few very wet districts cannot set the national profile. The shaded band marks June–September.




Figure E.5: Monsoon rainfall over time, local and upstream, by region

Notes: Yearly means across the estimation-sample districts of the two rainfall variables the regressions use: the district's own monsoon total and the river-distance-weighted monsoon total over its excl.-neighbour upstream catchment. Regions are assigned by state (Jammu & Kashmir; the seven Northeastern states; the rest), and the y-scales differ by region because the levels do.

E.2. Main-Result Robustness

Table E.1: The reduced form under alternative fixed effects, windows, and the downstream placebo

	Upstream ($t - 1$)	Uncertainty	Districts
<i>A. Fixed effects (published window)</i>			
District + year (baseline)	0.91	Conley 0.36	—
State $\times$ year	0.99	SE 0.39, $p = 0.010$	—
Region $\times$ year	1.20	$p = 0.002$	—
State linear trends	1.06	$p = 0.009$	—
<i>B. Sample windows</i>			
2000–2018 (crop window)	1.36	$t = 2.83$	—
2000–2025	1.68	$t = 3.88$	—
1989–2025	0.93	$t = 2.94$	213
<i>C. The downstream twin (annual rainfall)</i>			
Downstream alone, own sample	−1.71	SE 0.35	195
Downstream alone, common sample	−2.23	—	135
Downstream, three-way race	−2.33	—	135
Upstream in the same race	1.14	SE 0.43	135

Notes: Panel A varies the fixed effects on the published specification; state $\times$ year absorbs everything common to a state in a year, leaving districts of the same state draining different catchments. Panel B varies the window; the reduced form needs no crop data. Panel C reports the downstream placebo: the downstream series is built by the same procedure following the flow network in the opposite direction, the two treatments correlate 0.46 after fixed effects, and collinear readings of one river corridor split a common effect — the sign pattern, not the magnitude, is the informative object, and a misassigned upstream catchment would push the downstream term positive. Diagnostics retired from the prose: the lead of annual upstream rainfall is -0.01 ($t \approx 0$); upstream and local monsoon rainfall correlate 0.76 after fixed effects and the local controls are retained throughout; the effective number of independent shocks is 2,087 sub-basins, 37.4 PFAF-5 catchments and 7.20 macro basins, and a shock-level Borusyak–Hull–Jaravel transformation reproduces the point estimate with a first-stage F below 3, so nothing is concluded from it. Standard errors clustered on district unless stated.

Table E.2: Is the monsoon the right window?

	Monsoon	Crop-calendar alternatives		
	(1)	(2)	(3)	(4)
log Upstream Precip., monsoon ($t - 1$)	1.460*** (0.283)		1.547*** (0.301)	1.479*** (0.295)
log Upstream Precip., growing season ($t - 1$)		0.748*** (0.290)		
log Upstream Precip., non-monsoon window ($t - 1$)			-0.066 (0.119)	
Growing-season component orthogonal to the monsoon ($t - 1$)				-0.316 (0.270)
Observations	7,560	7,350	7,350	7,350
Number of Districts	210	210	210	210
Mean of dependent variable	1.00	1.01	1.01	1.01
Estimation	Poisson	Poisson	Poisson	Poisson
District & Year FE	✓	✓	✓	✓

Notes: Every column is the specification of Table 2, column 4, on 1990–2025; only the rainfall window changes. *Crop-calendar weights.* MIRCA2000 reports, for each district, the share of its cropped area under crops in each calendar month; these shares are time-invariant. Monthly upstream rainfall is multiplied by those shares and summed, so column 2 weights the whole year by when that district’s own agriculture actually happens rather than by a calendar common to all districts. *The non-monsoon term.* Columns 3 and 4 split the weights at the monsoon boundary and renormalise each half to sum to one, so the non-monsoon regressor carries *only* October–May rainfall, weighted by the share of the district’s cropping falling in those months — broadly the rabi season. That window straddles two calendar years, so it is dated by monsoon year: the term dated t runs from October of year t to May of year $t + 1$, and the lagged term covers the same eight months one year earlier. *Column 4.* The regressor is the residual from projecting the whole-year crop-calendar measure on monsoon rainfall within district and year — the crop-calendar content the monsoon window does not already carry. It is a generated regressor and its standard errors are not corrected for that first step. Current-period terms, weather controls and the spatial block are estimated throughout and not displayed. Conley spatial-HAC standard errors (150 km cutoff) in parentheses. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.




Figure E.6: The lagged upstream coefficient under alternative weightings

Notes: Each row re-estimates columns 1 and 4 of Table 2 with the treatment weights of the corresponding item in Section A.3; the specification, sample and controls are otherwise unchanged. Bars are 95 percent intervals on Conley spatial-HAC standard errors (150 km). The dashed line marks the paper’s estimate.




Figure E.7: How fast should weight decay with distance?

Notes: Each point re-estimates column 1 or column 4 of Table 2 with basin weights proportional to $(1 + d)^{-\gamma}$, times the basin's area where indicated, normalised within district; d is river or great-circle distance as in Figure E.6. $\gamma = 0$ is equal weighting and $\gamma = 1$ the paper's functional form. The band is the 95 percent Conley interval for river distance without area.




Figure E.8: The lagged upstream coefficient with rainfall averaged over coarser sub-basins

Notes: Each panel is one column of Table 2. The catchment, exclusion and river-distance weights are Table 2's, built from level-10 sub-basins. Each level-10 sub-basin's monsoon rainfall and temperature are replaced by the area-weighted average over the catchment's level-10 sub-basins that share its larger sub-basin at the level on the horizontal axis, so sub-basins removed by the exclusion never enter. The specification, sample and controls are unchanged. Level 10 is the paper's estimate, marked by the dashed line. Bars are 95 percent intervals on Conley spatial-HAC standard errors (150 km).

Table E.3: Upstream rainfall and conflict with the territory the census does not enumerate

	Baseline	Spatial controls			
	(1)	(2)	(3)	(4)	(5)
<i>A. As published: territory dropped</i>					
log Upstream Precip. ($t - 1$)	0.910** (0.359)	0.994*** (0.314)	1.351*** (0.351)	1.460*** (0.283)	1.242*** (0.277)
Observations	7,560	7,560	7,560	7,560	6,372
Number of Districts	210	210	210	210	177
Located events in sample	7,564	7,564	7,564	7,564	6,764
<i>B. Territory's events reassigned to the nearest Indian district</i>					
log Upstream Precip. ($t - 1$)	0.904** (0.356)	0.996*** (0.308)	1.343*** (0.359)	1.455*** (0.285)	1.237*** (0.279)
Observations	7,560	7,560	7,560	7,560	6,372
Number of Districts	210	210	210	210	177
Located events in sample	7,638	7,638	7,638	7,638	6,838
<i>C. Territory kept as a district</i>					
log Upstream Precip. ($t - 1$)	1.046*** (0.288)	0.965*** (0.276)	1.452*** (0.294)	1.470*** (0.260)	1.173*** (0.245)
Observations	7,596	7,596	7,596	7,596	6,372
Number of Districts	211	211	211	211	177
Located events in sample	8,218	8,218	8,218	8,218	6,764

Notes: Columns as in Table 2. The territory is the part of the former state of Jammu and Kashmir administered by Pakistan and China, which the 2001 census map draws as one unit but the census did not enumerate. Panel A is Table 2. Panel B assigns each of the territory's located events to the nearest Indian district; most go to districts with no upstream catchment under the exclusion rule, which is why few enter the sample. Panel C adds the territory as a district: its treatment is built as for every other district, from its upstream basins weighted by river distance under the exclusion rule, and it enters the neighbour matrix. It has no downstream catchment and so does not enter column 5, which moves only through the neighbour terms of the districts bordering it. Conley spatial-HAC standard errors (150 km cutoff) in parentheses. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.

Table E.4: Which conflicts respond, cut each way

	All conflict (1)	State-based (2)	One-sided (3)	Non-state (4)
<i>Panel A. By type of violence</i>				
log Upstream Precip. ($t - 1$)	1.460*** (0.283)	1.429*** (0.324)	1.063* (0.595)	2.491** (1.178)
log Upstream Precip. (t)	-0.086 (0.487)	-0.026 (0.511)	-0.378 (0.767)	0.555 (1.329)
log Local Precip. (t)	-0.088 (0.579)	0.049 (0.593)	-1.466 (0.910)	3.403** (1.679)
log Local Precip. ($t - 1$)	0.507 (1.025)	0.881 (1.238)	-0.936 (1.156)	5.595** (2.660)
1 SD of identifying variation (log points) equivalently, a wetter monsoon of (%)	0.12 12	0.12 12	0.12 12	0.11 12
Implied effect of a 1 SD wetter monsoon (%)	18	18	13	33
Observations	7,560	5,040	5,508	1,848
Districts	210	140	153	66
Mean of dependent variable	1.00	0.98	0.42	0.19
	Maoist (5)	Kashmir (6)	Northeast (7)	Communal (8)
<i>Panel B. By conflict system</i>				
log Upstream Precip. ($t - 1$)	1.206 (0.746)	0.722 (0.471)	1.758*** (0.595)	1.965 (1.274)
log Upstream Precip. (t)	-1.233* (0.723)	-0.403 (0.297)	0.624 (1.124)	1.107 (1.451)
log Local Precip. (t)	-1.091 (1.090)	-0.303 (0.428)	0.322 (0.742)	2.662 (2.430)
log Local Precip. ($t - 1$)	-1.545 (1.097)	0.909 (1.078)	-0.576 (0.394)	5.260 (3.225)
1 SD of identifying variation (log points) equivalently, a wetter monsoon of (%)	0.10 10	0.13 14	0.09 9	0.11 12
Implied effect of a 1 SD wetter monsoon (%)	13	10	17	25
Observations	3,255	1,548	1,190	1,334
Districts	93	43	34	58
Mean of dependent variable	0.95	1.36	1.45	0.20
Estimation	Poisson	Poisson	Poisson	Poisson
District & Year FE	✓	✓	✓	✓

Notes: Every cell is a separate Poisson regression on 1990–2025 under the specification of Table 2, column 4: district and year fixed effects, the full weather controls, and the spatial block. Only the outcome changes. Neighbours' conflict in the spatial block is built from the column's own outcome, so each column is conditioned on neighbouring violence of its own kind. Types are UCDP's type of violence, systems are assigned from the UCDP conflict name. Poisson drops districts recording no events of a given kind, which is why the district counts differ down the table. Conley spatial-HAC standard errors (150 km cutoff) in parentheses. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.

Table E.5: Where the upstream effect lands: own district against successive rings

	Own only (1)	+ Ring 1 (2)	+ Rings 1 and 2 (3)
<i>Panel A. Upstream monsoon rainfall</i>			
Own upstream	0.830*** (0.272)	0.765** (0.328)	0.832** (0.344)
Ring 1 (adjacent)		0.211 (0.538)	0.967 (0.597)
Ring 2 (adjacent to adjacent)			-1.421** (0.721)
Total	0.830*** (0.272)	0.975** (0.435)	0.379 (0.612)
Residual corr. with own	—	0.872	0.872, 0.816
<i>Panel B. Routed flood-stage days (CaMa)</i>			
Own upstream	0.019*** (0.007)	-0.004 (0.018)	0.011 (0.014)
Ring 1 (adjacent)		0.049 (0.036)	-0.008 (0.030)
Ring 2 (adjacent to adjacent)			0.074** (0.037)
Total	0.019*** (0.007)	0.045** (0.021)	0.077** (0.032)
Residual corr. with own	—	0.742	0.742, 0.458
Estimation	Poisson	Poisson	Poisson
District & Year FE	✓	✓	✓

Table E.6: What the spatial specification includes, and what happens if it does not

	log Upstream Precip. ($t - 1$)		Observations
<i>Panel A. Count outcome, fixed-effects Poisson</i>			
As estimated: $\lambda(Wy)_{t-1}$ and $\theta(WX)_t$	1.460***	(0.283)	7,560
+ own temporal lag γy_{t-1} (<i>inconsistent</i>)	1.410***	(0.423)	7,560
<i>Panel B. Linear outcome, $\log(1 + \text{count})$</i>			
Without γy_{t-1}	0.066**	(0.027)	15,300
+ own temporal lag γy_{t-1}	0.066***	(0.024)	15,300
+ contemporaneous spatial lag $\rho(Wy)_t$, by IV	0.142	(0.183)	15,300

Notes: Longest window, district and year fixed effects, Conley (150 km) standard errors in parentheses. The specification carries neighbours' lagged conflict and neighbours' weather but neither the contemporaneous spatial lag nor the district's own temporal lag, and this table gives the reason for each omission and shows what relaxing it does. *The own temporal lag.* A fixed-effects Poisson conditions on each district's own total count over the panel. $y_{i,t-1}$ is a component of that total and carries district i 's fixed effect mechanically, so conditioning no longer removes the incidental parameter and the estimator is inconsistent. The space-time lag $(Wy)_{i,t-1}$ is built from the *neighbours'* outcomes, whose fixed effects are absorbed by their own dummies and never enter i 's conditioning set; it is predetermined in the ordinary sense and needs only no-serial-correlation, not strict exogeneity. That distinction is what makes the combination principled rather than convenient. Row 2 imposes the term anyway: the elasticity is 1.410 against 1.460, so the omission conceals nothing. In the linear arm, where the term is admissible, adding it moves the coefficient from 0.0665 to 0.0659. Nickell bias is of order $1/T$ with $T = 36$ and is noted rather than solved, as in the source literature. *The contemporaneous spatial lag.* It is simultaneously determined — a district's violence enters its neighbours' regressor while theirs enters its own — so it requires either instruments or maximum likelihood. The conventional Kelejian–Pruha instrument, second-order neighbours' weather, has a first-stage F of 0.310; the resulting estimate is uninformative (last row) and $\hat{\rho}$ is -1.711. The instrument with real first-stage power on this panel — second-order neighbours' *upstream* exposure — is invalid, because catchments nest: it correlates 0.93 with a district's own treatment, so it reaches the outcome directly rather than only through the neighbours. The maximum-likelihood route of Harari and Ferrara (2018) is defined for a linear model of conflict incidence and has no standard counterpart for a fixed-effects Poisson count, which is the outcome here. Both conventional solutions to the simultaneity are therefore closed on the count margin, and the specification restricts to terms that need neither.

Table E.7: Spatial diffusion and the upstream treatment

	Incidence (LPM)				asinh(events)
	Contiguity (1)	River W (2)	Both (3)	Downstream (4)	River W (5)
Upstream rain ($t - 1$)	0.035* (0.019)	0.042** (0.019)	0.015 (0.036)	0.039** (0.019)	0.125*** (0.031)
Own conflict ($t - 1$)	0.170*** (0.026)	0.184*** (0.026)	0.169*** (0.026)	0.189*** (0.026)	0.427*** (0.037)
Neighbours' conflict ($t - 1$), contiguity	0.222*** (0.037)		0.194*** (0.041)		
Upstream districts' conflict ($t - 1$), river W		0.182*** (0.045)	0.061 (0.049)		0.144*** (0.053)
Downstream districts' rain (t)				-0.034* (0.019)	
<i>Panel B: event counts (Poisson), 176 districts</i>					
	Baseline	+ river conflict lag	+ contiguity too		
Upstream rain ($t - 1$)	0.861* (0.444)	1.134*** (0.384)	1.023*** (0.374)		
<i>Panel C: the flood arm — routed flood-stage days (Poisson), 152 districts</i>					
	Baseline	+ river lag	+ queen lag	+ upstr. flood days	
Flood-stage days (t)	0.024*** (0.006)	0.022*** (0.007)	0.023*** (0.007)	0.016* (0.008)	
Flood-stage days ($t - 1$)	0.019*** (0.007)	0.018*** (0.007)	0.017*** (0.006)	0.018*** (0.006)	
Upstream districts' flood days (t)				0.061* (0.035)	

Notes: Time-space recursive models with Durbin-type exogenous lags (Section E.2); 2000–2018; district and year fixed effects; standard errors clustered by district in parentheses. These are NOT dynamic spatial Durbin models in the sense of Harari and Ferrara (2018): the contemporaneous spatial lag $\rho(Wy)_t$ is absent throughout, for the reasons Table E.6 sets out. One consequence is interpretive and favourable — with no contemporaneous spatial lag there is no feedback loop, so these coefficients are effects directly and need no direct–indirect decomposition. Columns (1)–(4) are linear probability models of conflict incidence on all districts; column (5) uses asinh of the event count. Column (1) is the contiguity (queen) network and also holds neighbours' weather fixed; columns (2)–(5) use the directed river network, with (3) entering both conflict lags jointly and (4) the downstream direction. Panel B estimates the count model by fixed-effects Poisson with the own conflict lag excluded (strict exogeneity) on a frame without the upstream temperature controls, so its baseline sits below the published headline and comparisons are internal to the panel. Panel C repeats the exercise for the routed flood-stage days of the flood arm, local weather held throughout and the published no-controls form reproduced first; its final column adds upstream districts' own flood days, whose post-fixed-effects correlation with own days is 0.51 — a partial overlap with interpretable shares, not the near-collinear splits flagged in the text. ***, **, and * denote significance at the 1%, 5%, and 10% levels.




Figure E.9: The stability ladders: treatment coefficients as spatial controls enter, with 95 percent confidence intervals (district-clustered)

Notes: Top row: lagged upstream rainfall. Bottom row: contemporaneous routed flood-stage days. Left and centre columns are the linear models of Table E.7; the right column is the Poisson count model, where for rainfall the coefficient rises under diffusion controls and for flood days the final rung adds upstream districts' own flood days (the same-wave control, post-FE correlation 0.51). Poisson ladders omit the own-conflict lag throughout (strict exogeneity); the rainfall Poisson frame lacks the upstream temperature controls, so its baseline is not the published headline.




Figure E.10: Districts grouped by shared upstream catchment

Notes: Connected components of the district-by-district cosine similarity in upstream exposure, thresholded at 0.50. Exposure is the river-distance weight vector over HydroBASINS level-10 units that builds the treatment, so two districts are joined when more than 0.50 of their upstream water is shared, directly or through a chain of such links. Colours distinguish adjacent groups and carry no order. Districts drawn in flat grey have no upstream catchment in the construction. The partition is used in Section B as a clustering level for standard errors, not as an estimation sample.




Figure E.11: Randomization inference

Notes: Distribution of the lagged upstream elasticity across 5,000 within-state permutations of the district-catchment assignment; the red line is the estimate under the true assignment. Section B.4 describes the test.




Figure E.12: Four randomization tests of the same estimate

Notes: Section B.4 describes each test. Every panel re-estimates the specification of Table 2, column 4 on 1990–2025 under 5,000 permutations; the grey histogram is the null distribution, the vertical rule the observed statistic. Panel (a) permutes the district–catchment map within states, holding time fixed. Panels (b)–(d) permute the year labels of the whole cross-section jointly, holding the map fixed: one permutation per draw applied to every district, so the spatial dependence of the shock is preserved and only the shock–outcome alignment breaks. The outcome and the spatial terms built from it are never re-dated. Panel (c) restricts each permutation to stay within its decade, which preserves trend and relaxes the exchangeability the free permutation assumes. Panel (d) uses the score statistic implied by the Poisson likelihood. Permuting the upstream rainfall term alone rather than the whole forcing block gives $p = 0.0006$ on the studentized statistic; that narrower version leaves the treatment’s own temperature control and the local weather at their true dates, so the regressors are no longer mutually aligned.

Table E.8: Flooding at migrant origins and conflict at destinations

Origin definition	FloodMig (t)	Cluster SE	Perm. p (coef.)	Perm. p (t)
Origin state (D-02)	-0.059	(0.028)	0.109	0.150
Origin district (D-11)	0.017	(0.009)	0.532	0.244
Hybrid	0.016	(0.014)	0.613	0.526

Notes: Poisson fixed-effects, 2002–2018, the destination’s own routed flood days and local weather included in every column. FloodMig is the migration-share-weighted mean of origin flood days, with shares from the 2001 Census. D-02 gives origin states only; D-11 gives origin districts within a state. No definition produces an effect surviving the year-block permutation, and the district coefficient is a precise zero among destinations that never flood.




Figure E.13: Leave-one-out robustness, by district and by state

Notes: Each point re-estimates the specification of Table 2, column 4 on 1990–2025 with one district (panel a) or one state (panel b) removed; the red line is the full-sample estimate 1.46. Panel (a) plots the leave-it-out estimate against the dropped district’s own count of located events, so the vertical spread shows influence and the horizontal position shows concentration; the three labelled districts are the two extremes and the most violent. Panel (b) shows 95 percent district-clustered intervals, states ordered by the resulting estimate, with West Bengal — the most influential — marked. Districts are dropped from the frame before estimation, so the spatial block is rebuilt on the reduced sample rather than carried over.

Table E.9: Upstream rainfall partitioned by dam regulation

	Aggregate	Regulated	Free-flowing
log Upstream Precip. (t)	-0.851 (1.062)		
log Regulated Upstream Precip. (t)		-0.415 (0.872)	
log Free-Flowing Upstream Precip. (t)			-1.178 (1.026)
log Upstream Precip. ($t - 1$)	1.776*** (0.549)		
log Regulated Upstream Precip. ($t - 1$)		1.022*** (0.380)	
log Free-Flowing Upstream Precip. ($t - 1$)			2.402*** (0.884)
Observations	2052	2052	2052
Number of Districts	114	114	114
District & Year FE	✓	✓	✓

Notes: Poisson fixed-effects, UCDP event counts, 2000–2018, districts holding both a regulated and a free-flowing upstream component. Each column enters one component of the same upstream monsoon rainfall; local weather controls (current and lagged) included. The components share 80% of their within-district variation, so the contrast is estimated jointly in the text (total water plus composition) rather than by differencing these columns. Section VII shows the partition has no hydrological first stage and so does not identify buffering. Assam holds 2.7% of the identifying variation; excluding it, the free-minus-regulated gap falls to 0.32 against a placebo 95th percentile of 0.35 ($p = 0.11$). Placebo splits of the free-flowing basins on ten placement characteristics (and a propensity index) never reproduce the real gap, and stratified within-district permutation gives $p = 0.0002$ — direction, not magnitude, is the robust claim. Conley spatial-HAC standard errors (150 km cutoff) in parentheses, which is the paper’s headline estimator and the basis of the significance stars; Online Appendix Figure B.3 re-prices the headline coefficient under alternative variance estimators. ***, **, and * indicate statistical significance at the 1%, 5%, and 10% levels, respectively.

E.3. The Dam Arm

Table E.10: The dam arm’s first stage, and the routing counterfactual

	Free-flowing	Regulated	Gap (SE)
<i>A. Flood-stage days per log point of upstream monsoon rain</i>			
Routed days above the 2-year level	−3.26	13.03	−16.29 (3.20)
Districts		271	
<i>B. CaMa counterfactual: the same rain, dams on vs. off</i>			
Conflict on water that arrived		0.0184	
Conflict on water held back		0.0411 (0.0197)	
Districts		162	

Notes: Panel A regresses routed flood-stage days on log upstream monsoon rainfall separately for catchments with no dam (free-flowing) and with at least one (regulated), with district and year fixed effects; the gap is the difference. A buffering account requires the regulated catchment to deliver *less* flood per unit of rain; it delivers more, because dams sit on the largest rivers — the observational split has no usable first stage. Panel B holds the river fixed: CaMa-Flood is run twice on identical rainfall, with and without reservoir operation, and conflict is regressed on the water that arrived and the water held back, jointly. Attenuated delivery shows no differential conflict response.




Figure E.14: Permutation distribution for the free-minus-regulated gap, full sample

Notes: Distribution of the gap across 5,000 within-district redistributions of dam labels: within each district, which upstream sub-basins are labelled dam-regulated and which free-flowing is reassigned at random, keeping the number of each and holding the sub-basins, their weights and rainfall, and the outcomes fixed. The red line is the observed gap, 1.28 ($p = 0.0002$). Excluding Assam the observed gap falls to 0.32 and 537 of the same 5,000 placebo draws exceed it.

E.4. The Flood Arm




Figure E.15: Regional co-movement of conflict and routed flooding on the estimation sample, standardized within region (one axis; the two series have different units)

Notes: Located events per district (solid) and flood-stage days per district (dashed), 2001–2018. Jammu & Kashmir is omitted: its routed flood series is identically zero. The pooled flood coefficient survives every composition change but collapses under region-by-year effects while nearly all treatment variation survives, and alignment tests (leads, rotation randomization) cannot reject a coincidence of regional cycles; the flood channel’s evidentiary weight rests on the free-flowing versus dam-passing contrast, which shares this regional alignment by construction, and on the main text’s timing evidence.

References

- Abatzoglou, J. T., Dobrowski, S. Z., Parks, S. A., and Hegewisch, K. C. (2018). TerraClimate, a high-resolution global dataset of monthly climate and climatic water balance from 1958–2015. *Scientific Data*, 5:170191.
- Adão, R., Kolesár, M., and Morales, E. (2019). Shift-share designs: Theory and inference. *Quarterly Journal of Economics*, 134(4):1949–2010.
- Bates, P. D., Horritt, M. S., and Fewtrell, T. J. (2010). A simple inertial formulation of the shallow water equations for efficient two-dimensional flood inundation modelling. *Journal of Hydrology*, 387(1–2):33–45.
- Borusyak, K. and Hull, P. (2023). Nonrandom exposure to exogenous shocks. *Econometrica*, 91(6):2155–2185.
- Borusyak, K., Hull, P., and Jaravel, X. (2022). Quasi-experimental shift-share research designs. *Review of Economic Studies*, 89(1):181–213.
- Emerick, K., de Janvry, A., Sadoulet, E., and Dar, M. H. (2016). Technological innovations, downside risk, and the modernization of agriculture. *American Economic Review*, 106(6):1537–1561.
- Fetzer, T. (2020). Can Workfare Programs Moderate Conflict? Evidence from India. *Journal of the European Economic Association*, 18(6):3337–3375.
- Funk, C., Peterson, P., Harrison, L., et al. (2026). The Climate Hazards Center infrared precipitation with stations, version 3. *Scientific Data*, 13:718.
- Funk, C., Peterson, P., Landsfeld, M., et al. (2015). The climate hazards infrared precipitation with stations—a new environmental record for monitoring extremes. *Scientific Data*, 2:150066.
- Harari, M. and Ferrara, E. L. (2018). Conflict, Climate, and Cells: A Disaggregated Analysis. *The Review of Economics and Statistics*, 100(4):594–608.
- Harris, I., Osborn, T. J., Jones, P., and Lister, D. (2020). Version 4 of the CRU TS monthly high-resolution gridded multivariate climate dataset. *Scientific Data*, 7:109.
- Lehner, B. and Grill, G. (2013). Global river hydrography and network routing: Baseline data and new approaches to study the world’s large river systems. *Hydrological Processes*, 27(15):2171–2186.

- Muñoz-Sabater, J., Dutra, E., Agustí-Panareda, A., Albergel, C., Arduini, G., Balsamo, G., et al. (2021). ERA5-Land: A state-of-the-art global reanalysis dataset for land applications. *Earth System Science Data*, 13(9):4349–4383.
- Patel, D. (2024). Floods. Working paper, Harvard University.
- Schneider, U., Becker, A., Finger, P., Meyer-Christoffer, A., Ziese, M., and Rudolf, B. (2014). GPCC’s new land surface precipitation climatology based on quality-controlled in situ data and its role in quantifying the global water cycle. *Theoretical and Applied Climatology*, 115(1–2):15–40.
- Tellman, B., Sullivan, J. A., Kuhn, C., Kettner, A. J., Doyle, C. S., Brakenridge, G. R., Erickson, T. A., and Slayback, D. A. (2021). Satellite imaging reveals increased proportion of population exposed to floods. *Nature*, 596(7870):80–86.
- Yamazaki, D., de Almeida, G. A. M., and Bates, P. D. (2013). Improving computational efficiency in global river models by implementing the local inertial flow equation and a vector-based river network map. *Water Resources Research*, 49(11):7221–7235.
- Yamazaki, D., Ikeshima, D., Sosa, J., Bates, P. D., Allen, G. H., and Pavelsky, T. M. (2019). MERIT Hydro: A high-resolution global hydrography map based on latest topography dataset. *Water Resources Research*, 55(6):5053–5073.
- Yamazaki, D., Kanae, S., Kim, H., and Oki, T. (2011). A physically based description of floodplain inundation dynamics in a global river routing model. *Water Resources Research*, 47(4):W04501.
- Yamazaki, D., O’Loughlin, F., Trigg, M. A., Miller, Z. F., Pavelsky, T. M., and Bates, P. D. (2014a). Development of the global width database for large rivers. *Water Resources Research*, 50(4):3467–3480.
- Yamazaki, D., Sato, T., Kanae, S., Hirabayashi, Y., and Bates, P. D. (2014b). Regional flood dynamics in a bifurcating mega delta simulated in a global river model. *Geophysical Research Letters*, 41(9):3127–3135.